

HBase

★

RDBMS

HBASE

- Requires SQL
- Fixed schema
- Row oriented
- Not Scalable
- Static
- Slower retrieval of data
- It follows ACID prop.

- No SQL
- No fixed schema
- Column oriented
- Scalable
- Dynamic
- Faster retrieval
- It follows CAP (Consistency, Availability, Partition-tolerance) theorem
- Handles structured, semi-struct. & unstruct. data
- Can handle sparse data.
- Low latency

★ Limitations of hadoop

- Performs batch processing, sequential manner
- Needed to search entire dataset.
- Processing a huge dataset results into huge data sets, which is also processed sequentially

HBase.

- Open-source, horizontally scalable.
- Distributed column-oriented db build on top of Hadoop FS.
- Provides random real-time read/write access to data in Hadoop FS.

HDFS

- For storing files

- does not support fast individual record lookups
- High latency batch processing

- Sequential access of data

HBase

Storing in HDFS & processing.

Provides fast lookup for larger tables.
Low latency access to single row (Random access)

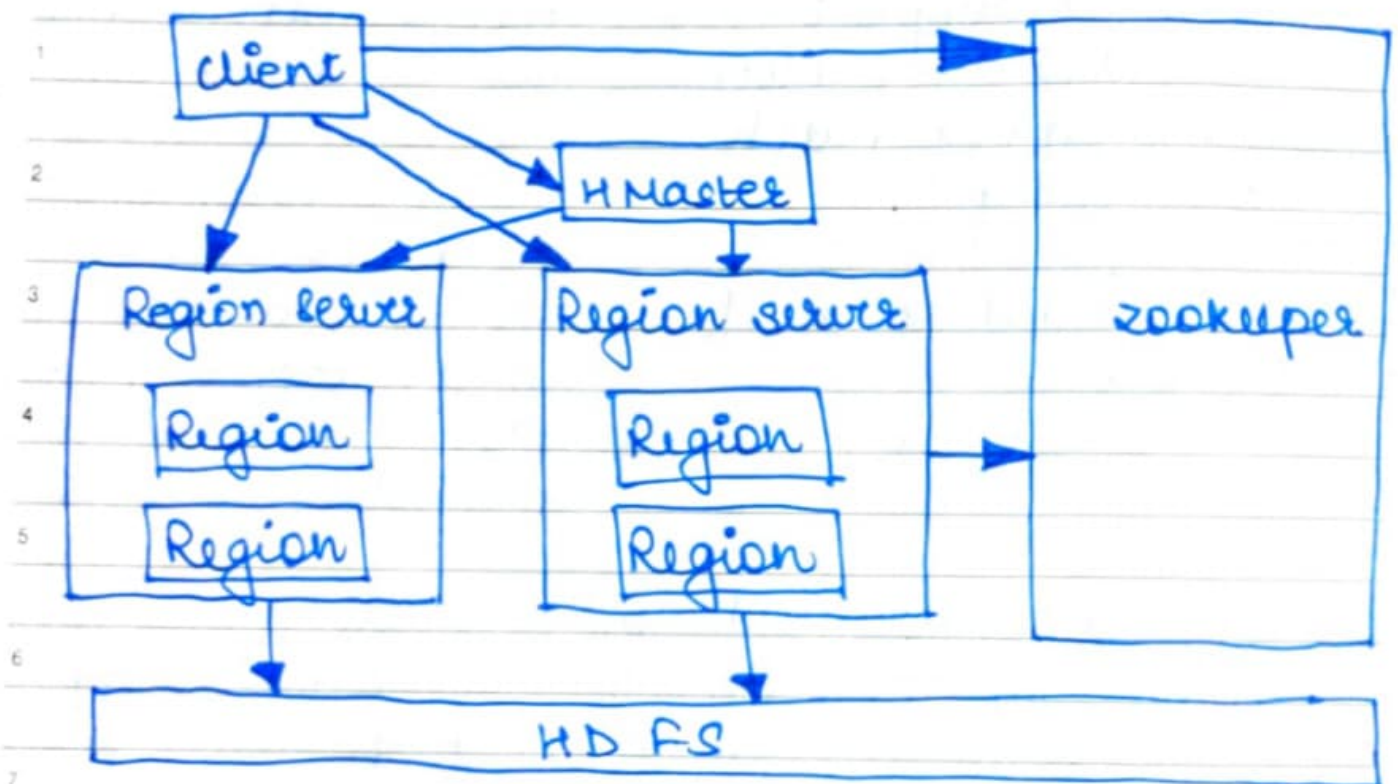
Uses hash tables & provides random access, & stores data in indexed HDFS for faster lookups.

• It is column-oriented db. and table are sorted by row. The table schema only ~~provides~~ defines column families, which are key-value pairs.

- Table is a collection of rows
- Row is a collection of column families
- Column family is a collection of columns
- Column is a collection of key-value pairs.

- Features of HBase
 - linearly scalable.
 - Automatic failure support.
 - Provides consistent reads & writes.
 - Integrated with Hadoop, both as source & destn.
 - Has easy JAVA API for client
 - Provides data replication across clusters.

* Architecture



1. HMaster :-

- Master Service of HBase
- Process in which regions are assigned to region server as well as DDL operations.
 - Table (create Table, remove, enable disable)
 - Column family (add column, modify column)
 - Region (move, assign)

- Manage region server instances present in cluster.
- Controlling Load balancing, failover etc.
- Acts as an interface for DDL operations.

2. Zookeeper :-

- Maintains server state, coordination service.
- Provides services like maintaining configuration information, naming, providing distributed synchronization, server failure notification etc.
- Clients communicate with region servers via zookeeper

3. Region Server

- HBase Tables are divided horizontally by row key range into Regions.
- Regions are the basic building elements of HBase cluster that consists of distribution of tables & are comprised of column families.
- Runs on HDFS DN.
- Regions are responsible for handling, managing executing as well as reads & writes HBase operations on that set of region.
- Default size of region → 256 MB.

→ HMaster → Region server

- Hosting & managing regions

- Splitting regions automatically

- Handling read & writes requests.

- Communicates with client directly

Region → building element of HBase consist of distributed tables & comprised of column families. contain multiple store, one for each column family. Has memstore, n file.

05

SATURDAY

FEBRUARY

★ MetaTable. (stored on Region Servers)

→ It holds the location of the Regions in HBase cluster.

→ It keeps list of all Regions in the system

→ Structure of Metatable:

1. Key: region start key, region id.

2. Values: Region Server.

→ Handled by Zookeeper.

- It is like a binary tree.

★ HBase Write Steps

① First step: Is to write the data to the write-ahead log, while the client issues a put request:

→ To the end of WAL file, all the edits are appended which is stored on disk.

→ In case a server crashes, WAL is used to recover not-yet-persisted data

② As soon as data is written to WAL, it is stored in memstore. (256mb written in H Files)
in chunks 100mb

★ Region Server Components

1. WAL (Write-ahead log)

→ Stores new data to permanent storage.

→ Also used for recovery.

2. Block Cache

→ It is the read cache.

FEB
2022

M T W T F S S M T W T F S S M T W T F S S M T W T F S S
• 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 • • • • •

read

- stores frequent data in memory & removes least read data when full.

3. MemStore

- write cache, stores new data which has not been written to disk.
- sorts before writing.
- One MemStore per column family

4. HFiles

- These files stores the rows as sorted key value on disk.

* Compaction

- Combining of HFiles to reduce the storage & reduce the no. of disk seeks needed for read.

1. Minor : Picks smaller HFiles & combines them to bigger Hfile. It performs merge sort.

2. Major : The bigger Hfile the same column families are placed together.

- Scheduled during low peak load timings to avoid congestion.

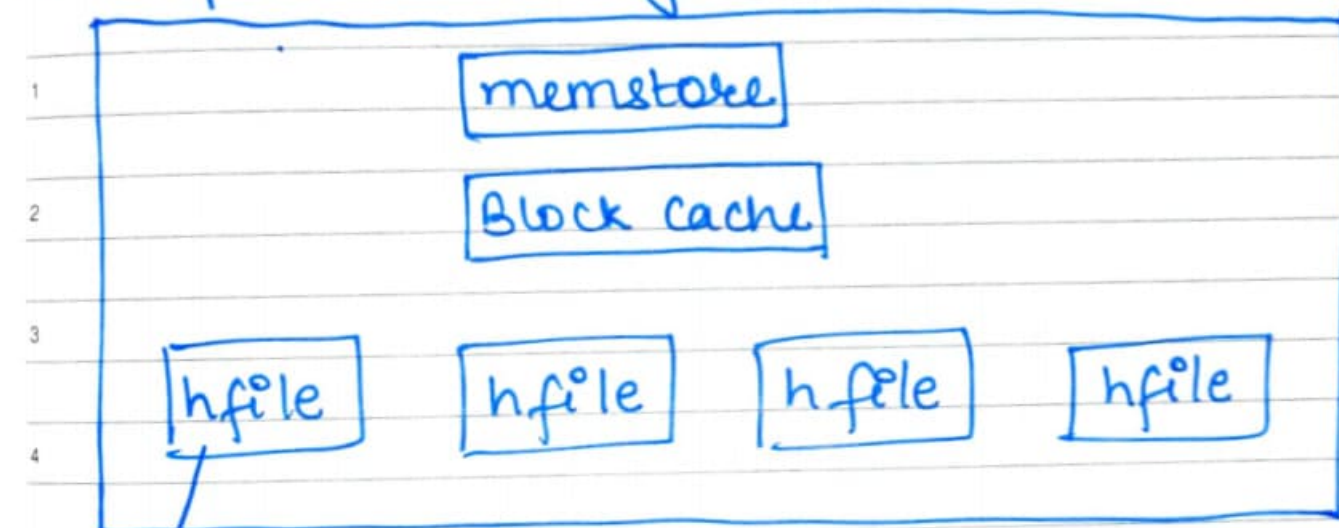
* HBASE Architecture

• Zookeeper → Hmaster → Region Server
→ Regions → Column Family Store.

• Region

→ Actual data is splitted in regions
→ size (default) of region → 256 mb. (for optimal performance).

• Components in Region



size is in
kbs

storing column families

Write operation

→ data is sent to WAL & then to Memstore → after memory is full of memstore it flushes data into hfiles.

Client

Region Server

HLOG/WAL

Region

memstore

Block cache

hfile

hfile

hfile

hfile

column family storage

Before writing data on disk it is written in buffer which is memstore (assume 100 mb)

M T W T F S S M T W T F S S M T W T F S S M T W T F S S M T W T F S S M T W T F S S
• • • • 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 2022

05

DAY 156-209 WEEK 23

SUNDAY

2022

JUNE

→ Read Operation

client → Metatable on Region Servers &
→ It gives location of region where
table is stored → memstore, Block
cache, hfiles are scanned for data. →
given to client.

* Commands

- status - Provides the status of nBase, no. of servers
- version - Provides the version of nBase
- table_help - Provides help for table-reference cmds.
- whoami - Provides information abt the user.

* Data Definition Language

- create, list, disable, is_disabled, enable, is_enabled, describe, alter, exists, drop, drop_all,

* bml

- put₄ - puts a cell value at a specified column is a specified row in a particular table.
- get - fetch contents of row or a cell.
- delete₅ - delete a cell value in a table
- delete₆ all - delete all the cells in a given row.
- scan, count, truncate,

2022

FEBRUARY

DAY 039-326 WEEK 07

TUESDAY

08

eg. ~~1~~

Row Key	Data
cutting	info: { 'height': '9ft', 'state': 'CA' } roles: { 'ASF': 'Director', 'Hadoop': 'Founder' }
tipcon	info: { 'height': '5ft7', 'state': 'CA' } roles: { 'Hadoop': 'Committee' @ ts = 2010, 'Hadoop': 'PMC' @ ts = 2011, 'Hive': 'Contributor' }

Creates sp. H base
tables of column
Families.

info Column Family

Row Key	Columnkey	Timestamp	Cell value
cutting	info: height	1273516197868	9ft
cutting	info: state	1043871824184	CA
tipcon	info: height	1273878447049	5ft 7
tipcon	info: state	1273616297446	CA

Roles Column Family

Row Key	Columnkey	Timestamp	Cell value
cutting	roles: ASF	1273871823022	Director
cutting	roles: Hadoop	1183746289103	Founder
tipcon	roles: Hadoop	...	PMC
tipcon	roles: Hadoop	...	Committee
tipcon	roles: Hive	...	Contributor

Sorted on disk by Row key, column key,
desc. timestamp.