

→ Combiner

- Called local reducer
- Runs on same node where mapper is running
- Combines data

```
mapper1
A,10
A,20
A,50
B,20
B,30

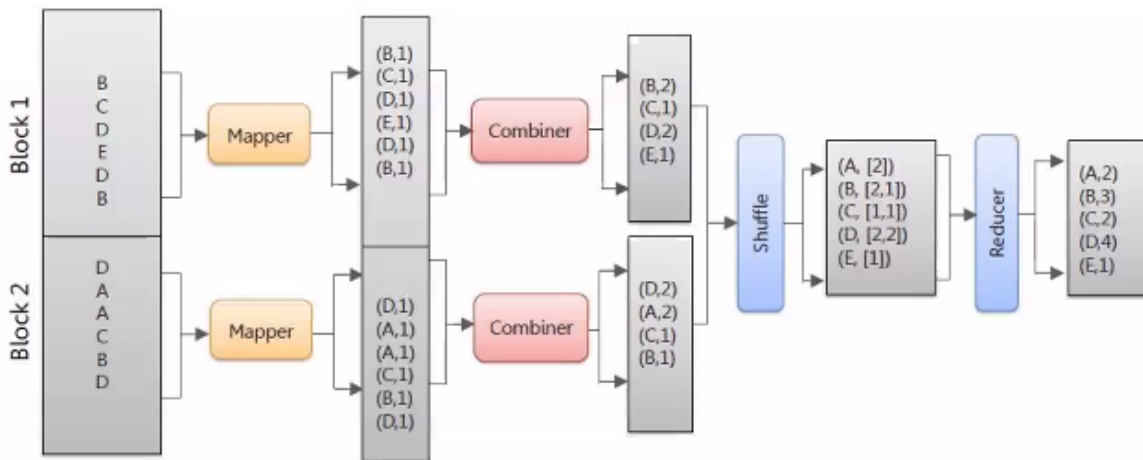
mapper2
A,5
A,20
B,10
B,60

A, [10,20,50,5,20]
B, [20,30,10,60]

Reducer - addition

A,105
B,120

without comb
AA      42061448400
```



```
Untitled - Notepad
File Edit Format View Help
after combiner
AA      42061448400

m1
A,76
A,70

m2
A,60

A, [76,70,60]

sum = 206
cnt = 3

avg = 206/3 = 68.66

A,68.66

-----
m1
A,76
A,70

c1

m2
```

→

```
import java.io.*;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.fs.*;
import org.apache.hadoop.mapreduce.lib.input.*;
import org.apache.hadoop.mapreduce.lib.output.*;
public class StockVolumeWithCombiner {

    public static class MapClass extends Mapper<LongWritable,Text,Text,LongWritable>
    {
        public void map(LongWritable key, Text value, Context context)
        {
            try{
                String[] str = value.toString().split(",");
                long vol = Long.parseLong(str[7]);
                context.write(new Text(str[1]),new LongWritable(vol));
            }
            catch(Exception e)
```

```

        {
            System.out.println(e.getMessage());
        }
    }
}

public static class ReduceClass extends Reducer<Text,LongWritable,Text,LongWritable>
{
    private LongWritable result = new LongWritable();

    public void reduce(Text key, Iterable<LongWritable> values,Context context) throws
IOException, InterruptedException {
        long sum = 0;

        for (LongWritable val : values)
        {
            sum += val.get();
        }

        result.set(sum);
        context.write(key, result);
        //context.write(key, new LongWritable(sum));
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    //conf.set("name", "value")
    //conf.set("mapreduce.input.fileinputformat.split.minsize", "134217728");
    Job job = Job.getInstance(conf, "Volume Count");
    job.setJarByClass(SockVolume.class);
    job.setMapperClass(MapClass.class);
    //job.setCombinerClass(ReduceClass.class);
    job.setCombinerClass(ReduceClass.class); //enables combiner class to run at mapper
    job.setReducerClass(ReduceClass.class);
    //
    job.setNumReduceTasks(1); // means it will run one reducer only
    job.setNumReduceTasks(0); // means no reducer will run
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(LongWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

```

[bigdatalab456422@ip-10-1-1-204 ~]$ jar tvf myjar.jar
25 Sat May 20 15:05:20 UTC 2023 META-INF/MANIFEST.MF

```

```

2459 Fri May 19 16:03:30 UTC 2023 AllTimeHigh$MapClass.class
2392 Fri May 19 16:03:30 UTC 2023 AllTimeHigh$ReduceClass.class
1722 Fri May 19 16:03:30 UTC 2023 AllTimeHigh.class
2475 Fri May 19 16:53:46 UTC 2023 AvgClosingPrice$MapClass.class
2454 Fri May 19 16:53:46 UTC 2023 AvgClosingPrice$ReduceClass.class
1732 Fri May 19 16:53:46 UTC 2023 AvgClosingPrice.class
2337 Fri May 19 17:41:44 UTC 2023 WordCount$IntSumReducer.class
2461 Fri May 19 17:41:44 UTC 2023 WordCount$TokenizerMapper.class
1790 Fri May 19 17:41:44 UTC 2023 WordCount.class
2454 Fri May 19 15:53:50 UTC 2023 AllTimeLow$MapClass.class
2388 Fri May 19 15:53:50 UTC 2023 AllTimeLow$ReduceClass.class
1734 Fri May 19 15:53:50 UTC 2023 AllTimeLow.class
2408 Thu May 18 17:48:56 UTC 2023 StockVolume$MapClass.class
2349 Thu May 18 17:48:56 UTC 2023 StockVolume$ReduceClass.class
1697 Thu May 18 17:48:56 UTC 2023 StockVolume.class
2456 Sat May 20 15:04:24 UTC 2023
StockVolumeWithCombiner$MapClass.class
2397 Sat May 20 15:04:24 UTC 2023
StockVolumeWithCombiner$ReduceClass.class
1813 Sat May 20 15:04:24 UTC 2023 StockVolumeWithCombiner.class

```

```

[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop jar myjar.jar
StockVolumeWithCombiner training/NYSE.csv training/out7
WARNING: Use "yarn jar" to launch YARN applications.
23/05/20 09:39:36 INFO client.RMPProxy: Connecting to ResourceManager
at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
23/05/20 09:39:37 WARN mapreduce.JobResourceUploader: Hadoop
command-line option parsing not performed. Implement the Tool
interface and execute your application with ToolRunner to remedy this.
23/05/20 09:39:37 INFO mapreduce.JobResourceUploader: Disabling
Erasure Coding for path:
/user/bigdatalab456422/.staging/job_1684298513961_0758
23/05/20 09:39:37 INFO input.FileInputFormat: Total input files to
process : 1
23/05/20 09:39:37 INFO mapreduce.JobSubmitter: number of splits:1
23/05/20 09:39:37 INFO Configuration.deprecation:
yarn.resourcemanager.system-metrics-publisher.enabled is deprecated.
Instead, use yarn.system-metrics-publisher.enabled
23/05/20 09:39:37 INFO mapreduce.JobSubmitter: Submitting tokens for
job: job_1684298513961_0758

```

```
23/05/20 09:39:37 INFO mapreduce.JobSubmitter: Executing with tokens:
[]
23/05/20 09:39:37 INFO conf.Configuration: resource-types.xml not
found
23/05/20 09:39:37 INFO resource.ResourceUtils: Unable to find
'resource-types.xml'.
23/05/20 09:39:37 INFO impl.YarnClientImpl: Submitted application
application_1684298513961_0758
23/05/20 09:39:37 INFO mapreduce.Job: The url to track the job:
http://ip-10-1-1-204.ap-south-1.compute.internal:6066/proxy/applicati
on\_1684298513961\_0758/
23/05/20 09:39:37 INFO mapreduce.Job: Running job:
job_1684298513961_0758
23/05/20 09:39:47 INFO mapreduce.Job: Job job_1684298513961_0758
running in uber mode : false
23/05/20 09:39:47 INFO mapreduce.Job:  map 0% reduce 0%
23/05/20 09:39:53 INFO mapreduce.Job:  map 100% reduce 0%
23/05/20 09:39:53 INFO mapreduce.Job: Job job_1684298513961_0758
completed successfully
23/05/20 09:39:53 INFO mapreduce.Job: Counters: 33
    File System Counters
        FILE: Number of bytes read=0
        FILE: Number of bytes written=222643
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=40990986
        HDFS: Number of bytes written=7842509
        HDFS: Number of read operations=7
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
        HDFS: Number of bytes read erasure-coded=0
    Job Counters
        Launched map tasks=1
        Data-local map tasks=1
        Total time spent by all maps in occupied slots
(ms)=4110
        Total time spent by all reduces in occupied slots
(ms)=0
        Total time spent by all map tasks (ms)=4110
        Total vcore-milliseconds taken by all map tasks=4110
        Total megabyte-milliseconds taken by all map
tasks=4208640
    Map-Reduce Framework
        Map input records=735026
```

```
Map output records=735026
Input split bytes=124
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=68
CPU time spent (ms)=3180
Physical memory (bytes) snapshot=370642944
Virtual memory (bytes) snapshot=2587979776
Total committed heap usage (bytes)=468713472
Peak Map Physical memory (bytes)=370642944
Peak Map Virtual memory (bytes)=2587979776

File Input Format Counters
  Bytes Read=40990862
File Output Format Counters
  Bytes Written=7842509
```

→ customized combiner

```
-----
customized combiner

m1
A, 76
A, 70
..
c1
A, 146, 2
A, 60

m2
A, 60

c2
A, (60, 1)
A, [(146, 2), (60, 1)]

sum = 146+60
cnt=2+1

avg = 206/3

A, 68.66
```

Handwritten annotations in red:

- Arrows pointing from `m1` to `A, 76` and `A, 70` with the text "Text, Text".
- An arrow pointing from `c1` to `A, 146, 2` with the text "sum".
- An arrow pointing from `A, 146, 2` to `A, 60` with the text "Text, Text".
- A box containing `m` and `c` with an arrow pointing to `sum`.

→ customized combiner: find out average score for students of each city (check if Pune city has average score 68.6666...)

#

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -mkdir student
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -put student1 student2
student
```

#

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop jar myjar.jar CityAvg2
student training/out8
WARNING: Use "yarn jar" to launch YARN applications.
23/05/20 10:25:44 INFO client.RMProxy: Connecting to ResourceManager
at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
23/05/20 10:25:45 WARN mapreduce.JobResourceUploader: Hadoop
command-line option parsing not performed. Implement the Tool
interface and execute your application with ToolRunner to remedy this.
23/05/20 10:25:45 INFO mapreduce.JobResourceUploader: Disabling
Erasure Coding for path:
/user/bigdatalab456422/.staging/job_1684298513961_0784
23/05/20 10:25:45 INFO input.FileInputFormat: Total input files to
process : 2
23/05/20 10:25:46 INFO mapreduce.JobSubmitter: number of splits:2
23/05/20 10:25:46 INFO Configuration.deprecation:
yarn.resourcemanager.system-metrics-publisher.enabled is deprecated.
Instead, use yarn.system-metrics-publisher.enabled
23/05/20 10:25:46 INFO mapreduce.JobSubmitter: Submitting tokens for
job: job_1684298513961_0784
23/05/20 10:25:46 INFO mapreduce.JobSubmitter: Executing with tokens:
[]
23/05/20 10:25:46 INFO conf.Configuration: resource-types.xml not
found
23/05/20 10:25:46 INFO resource.ResourceUtils: Unable to find
'resource-types.xml'.
23/05/20 10:25:46 INFO impl.YarnClientImpl: Submitted application
application_1684298513961_0784
23/05/20 10:25:46 INFO mapreduce.Job: The url to track the job:
http://ip-10-1-1-204.ap-south-1.compute.internal:6066/proxy/applicati
on\_1684298513961\_0784/
23/05/20 10:25:46 INFO mapreduce.Job: Running job:
job_1684298513961_0784
23/05/20 10:26:09 INFO mapreduce.Job: Job job_1684298513961_0784
running in uber mode : false
23/05/20 10:26:09 INFO mapreduce.Job: map 0% reduce 0%
```

23/05/20 10:26:17 INFO mapreduce.Job: map 100% reduce 0%
23/05/20 10:26:31 INFO mapreduce.Job: map 100% reduce 100%
23/05/20 10:26:33 INFO mapreduce.Job: Job job_1684298513961_0784
completed successfully

23/05/20 10:26:34 INFO mapreduce.Job: Counters: 54

File System Counters

FILE: Number of bytes read=148
FILE: Number of bytes written=670328
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=493
HDFS: Number of bytes written=101
HDFS: Number of read operations=11
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=2
Launched reduce tasks=1
Data-local map tasks=2
Total time spent by all maps in occupied slots
(ms)=10570
Total time spent by all reduces in occupied slots
(ms)=12278
Total time spent by all map tasks (ms)=10570
Total time spent by all reduce tasks (ms)=12278
Total vcore-milliseconds taken by all map tasks=10570
Total vcore-milliseconds taken by all reduce
tasks=12278
Total megabyte-milliseconds taken by all map
tasks=10823680
Total megabyte-milliseconds taken by all reduce
tasks=12572672

Map-Reduce Framework

Map input records=12
Map output records=12
Map output bytes=122
Map output materialized bytes=179
Input split bytes=246
Combine input records=12
Combine output records=10
Reduce input groups=7
Reduce shuffle bytes=179
Reduce input records=10


```
Reduce output records=7
Spilled Records=20
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=571
CPU time spent (ms)=3370
Physical memory (bytes) snapshot=1306103808
Virtual memory (bytes) snapshot=7778222080
Total committed heap usage (bytes)=1710227456
Peak Map Physical memory (bytes)=526311424
Peak Map Virtual memory (bytes)=2590330880
Peak Reduce Physical memory (bytes)=256356352
Peak Reduce Virtual memory (bytes)=2599239680

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=247
File Output Format Counters
  Bytes Written=101
```

[↩ Back](#)[🔗 Edit file](#)[🔄 Refresh](#)[📄 View as binary](#)[📄 Download](#)

Last modified
05/20/2023 3:56 PM

[🏠 Home](#)

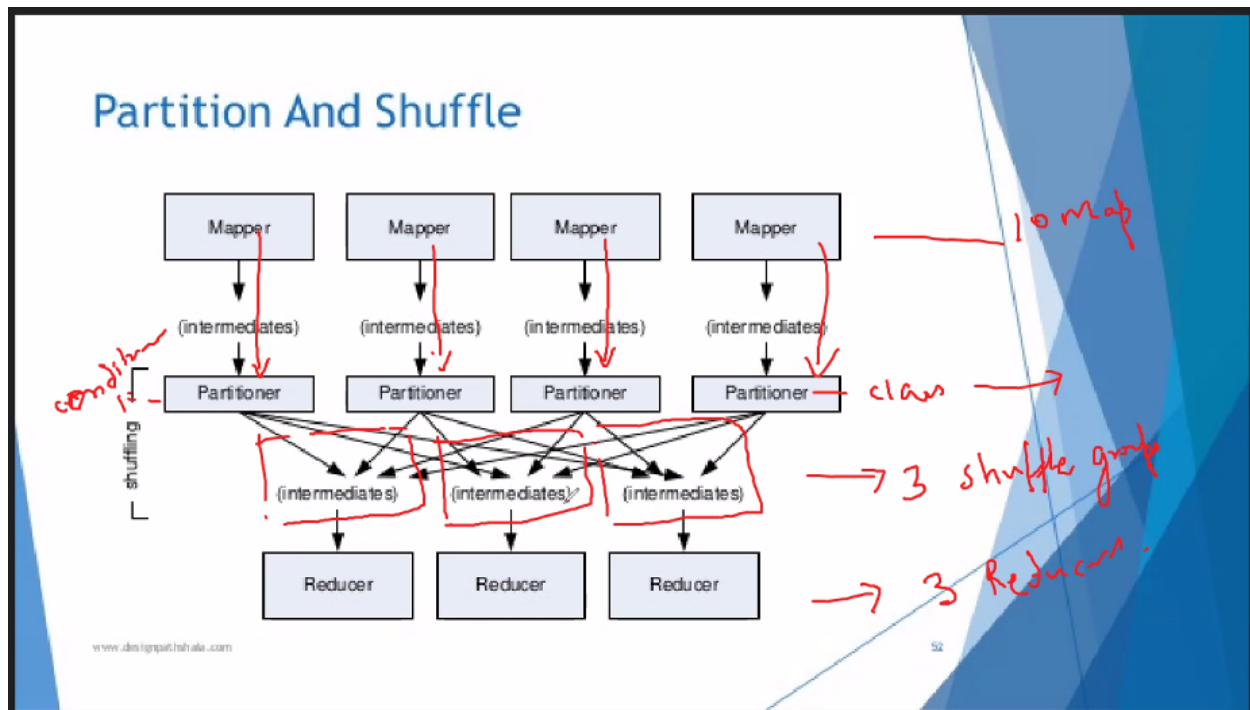
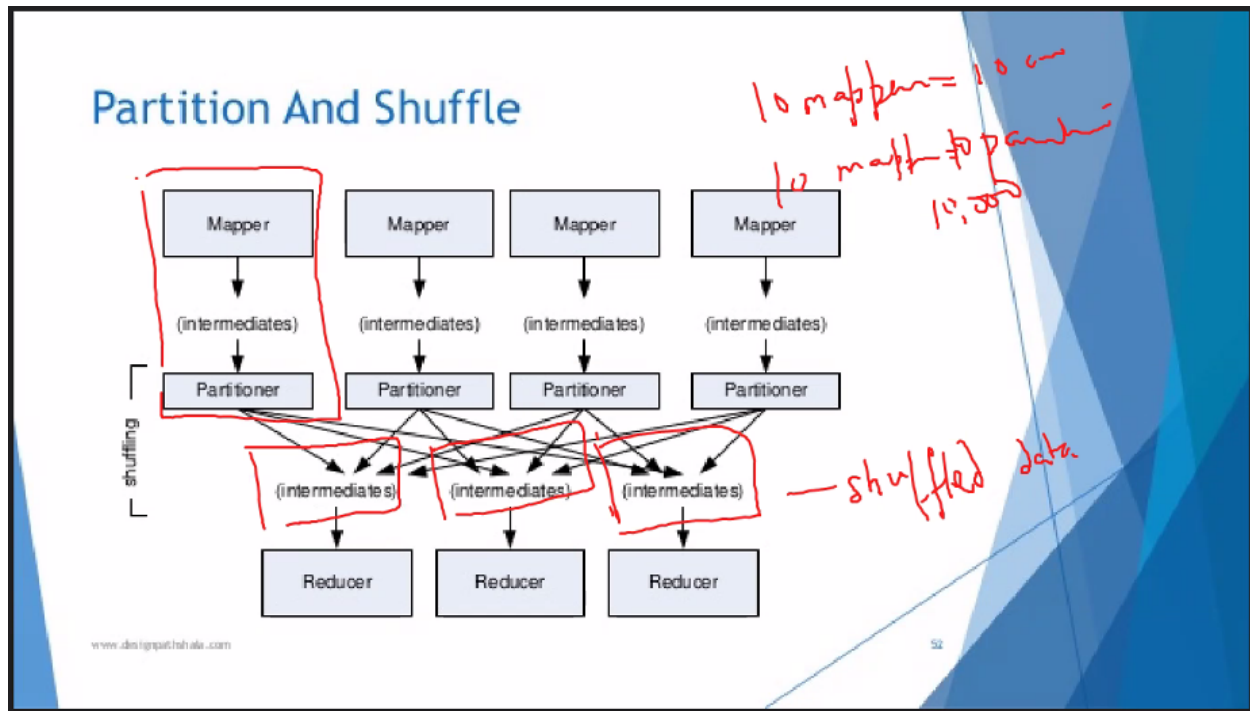
Page to of 1 [⏪](#) [⏴](#) [⏵](#) [⏩](#)

/ [user](#) / [bigdatalab456422](#) / [training](#) / [out8](#) / **[part-r-00000](#)**

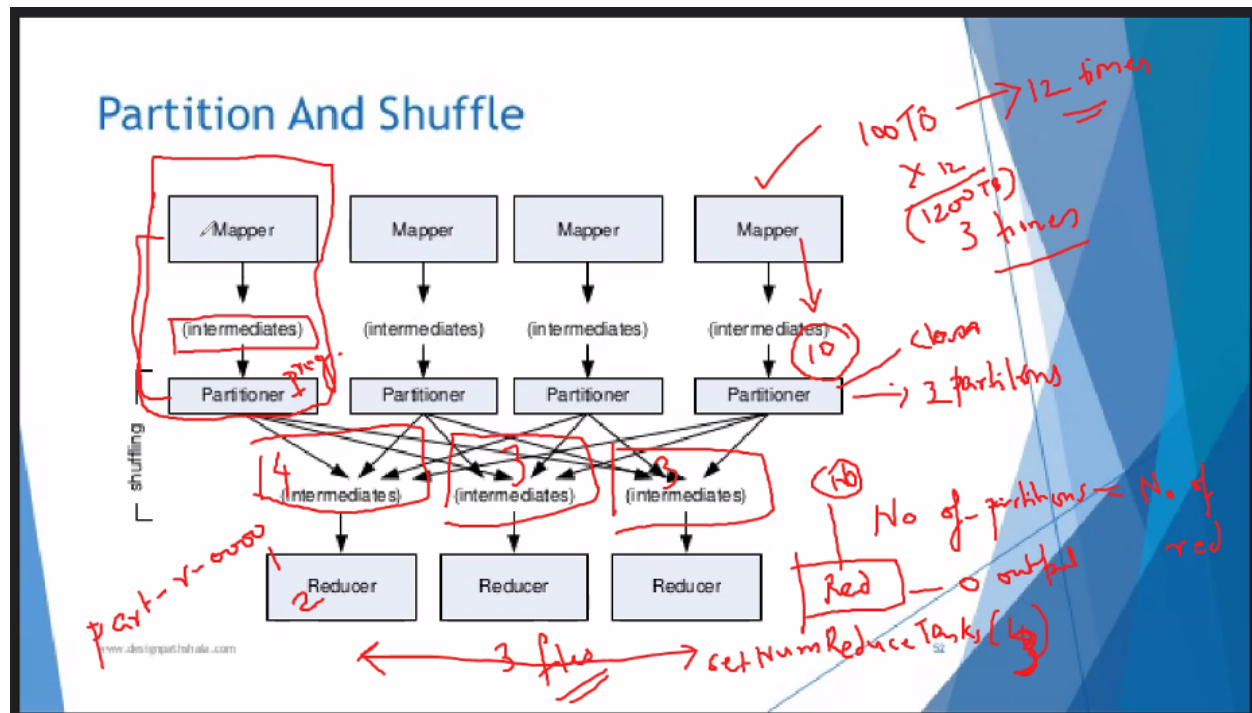
```
Bangalore,85.0
Chennai,90.0
Delhi,83.0
Hyderabad,80.0
Jaipur,75.0
Mumbai,86.0
Pune,68.66666666666667
```

→ Partitioner

- Called local reducer
- Runs on same node where mapper is running



Partition And Shuffle



```
emp.csv - Notepad
File Edit Format View Help
1201,gopal,45,Male,50000
1202,manisha,40,Female,51000
1203,khaleel,34,Male,30000
1204,prasanth,30,Male,31000
1205,kiran,20,Male,40000
1206,laxmi,25,Female,35000
1207,bhavya,20,Female,15000
1208,reshma,19,Female,14000
1209,kranthi,22,Male,22000
1210,Satish,24,Male,25000
1211,Krishna,25,Male,26000
1212,Arshad,28,Male,20000
1213,lavanya,18,Female,8000

prob statement : find max salary paid to each gender in 3 diff age groups

1) age<=20
2) age > and <=30
3) age > 30

mapper code
gender, (age,sal)
```

```
emp.csv - Notepad
File Edit Format View Help

1) age<=20
2) age > and <=30
3) age > 30

mapper code
gender, (age,sal)

1201,gopal,45,Male,50000
1202,manisha,40,Female,51000
1203,khaleel,34,Male,30000
1204,prasanth,30,Male,31000

m1
```

```
emp.csv - Notepad
File Edit Format View Help

1204,prasanth,30,Male,31000

m1
Male, (45,50000)
Female, (40,51000)
Male, (34,30000)
Male, (30,31000)

p-0
no records

p-1
Male, (30,31000)

p-2
Male, (45,50000)
Female, (40,51000)
Male, (34,30000)

r-2
Female, [(40,51000)]
Male, [(45,50000),(34,30000)]

F - 51K
M - 50K
```

→ find max salary paid to each gender in 3 diff age groups

- 1) age<=20
- 2) age > 20 and age <= 30
- 3) age > 30

#

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -put emp.csv training
```

#

```
[bigdatalab456422@ip-10-1-1-204 ~]$ jar tvf myjar.jar
 25 Sat May 20 17:36:10 UTC 2023 META-INF/MANIFEST.MF
2459 Fri May 19 16:03:30 UTC 2023 AllTimeHigh$MapClass.class
2392 Fri May 19 16:03:30 UTC 2023 AllTimeHigh$ReduceClass.class
1722 Fri May 19 16:03:30 UTC 2023 AllTimeHigh.class
2475 Fri May 19 16:53:46 UTC 2023 AvgClosingPrice$MapClass.class
```

```

2454 Fri May 19 16:53:46 UTC 2023 AvgClosingPrice$ReduceClass.class
1732 Fri May 19 16:53:46 UTC 2023 AvgClosingPrice.class
2337 Fri May 19 17:41:44 UTC 2023 WordCount$IntSumReducer.class
2461 Fri May 19 17:41:44 UTC 2023 WordCount$TokenizerMapper.class
1790 Fri May 19 17:41:44 UTC 2023 WordCount.class
2454 Fri May 19 15:53:50 UTC 2023 AllTimeLow$MapClass.class
2388 Fri May 19 15:53:50 UTC 2023 AllTimeLow$ReduceClass.class
1734 Fri May 19 15:53:50 UTC 2023 AllTimeLow.class
1242 Sat May 20 17:28:50 UTC 2023
MyPartitioner$CaderPartitioner.class
2365 Sat May 20 17:28:50 UTC 2023 MyPartitioner$MapClass.class
2905 Sat May 20 17:28:50 UTC 2023 MyPartitioner$ReduceClass.class
2632 Sat May 20 17:28:50 UTC 2023 MyPartitioner.class
2408 Thu May 18 17:48:56 UTC 2023 StockVolume$MapClass.class
2349 Thu May 18 17:48:56 UTC 2023 StockVolume$ReduceClass.class
1697 Thu May 18 17:48:56 UTC 2023 StockVolume.class
2648 Sat May 20 15:42:48 UTC 2023 CityAvg2$CityCombineClass.class
2269 Sat May 20 15:42:48 UTC 2023 CityAvg2$CityMapClass.class
2639 Sat May 20 15:42:48 UTC 2023 CityAvg2$CityReduceClass.class
2034 Sat May 20 15:42:48 UTC 2023 CityAvg2.class
2456 Sat May 20 15:04:24 UTC 2023
StockVolumeWithCombiner$MapClass.class
2397 Sat May 20 15:04:24 UTC 2023
StockVolumeWithCombiner$ReduceClass.class
1813 Sat May 20 15:04:24 UTC 2023 StockVolumeWithCombiner.class

```

#

```

[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop jar myjar.jar
MyPartitioner training/emp.csv training/out9
WARNING: Use "yarn jar" to launch YARN applications.
23/05/20 12:09:40 INFO client.RMPProxy: Connecting to ResourceManager
at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
23/05/20 12:09:40 WARN mapreduce.JobResourceUploader: Hadoop
command-line option parsing not performed. Implement the Tool
interface and execute your application with T
oolRunner to remedy this.
23/05/20 12:09:40 INFO mapreduce.JobResourceUploader: Disabling
Erasure Coding for path:
/user/bigdatalab456422/.staging/job_1684298513961_0866
23/05/20 12:09:41 INFO input.FileInputFormat: Total input files to
process : 1
23/05/20 12:09:41 INFO mapreduce.JobSubmitter: number of splits:1

```

```
23/05/20 12:09:41 INFO Configuration.deprecation:
yarn.resourcemanager.system-metrics-publisher.enabled is deprecated.
Instead, use yarn.system-metrics-publisher.enabled
23/05/20 12:09:41 INFO mapreduce.JobSubmitter: Submitting tokens for
job: job_1684298513961_0866
23/05/20 12:09:41 INFO mapreduce.JobSubmitter: Executing with tokens:
[]
23/05/20 12:09:41 INFO conf.Configuration: resource-types.xml not
found
23/05/20 12:09:41 INFO resource.ResourceUtils: Unable to find
'resource-types.xml'.
23/05/20 12:09:42 INFO impl.YarnClientImpl: Submitted application
application_1684298513961_0866
23/05/20 12:09:42 INFO mapreduce.Job: The url to track the job:
http://ip-10-1-1-204.ap-south-1.compute.internal:6066/proxy/applicati
on\_1684298513961\_0866/
23/05/20 12:09:42 INFO mapreduce.Job: Running job:
job_1684298513961_0866
23/05/20 12:09:57 INFO mapreduce.Job: Job job_1684298513961_0866
running in uber mode : false
23/05/20 12:09:57 INFO mapreduce.Job:  map 0% reduce 0%
23/05/20 12:10:04 INFO mapreduce.Job:  map 100% reduce 0%
23/05/20 12:10:13 INFO mapreduce.Job:  map 100% reduce 67%
23/05/20 12:10:16 INFO mapreduce.Job:  map 100% reduce 100%
23/05/20 12:10:17 INFO mapreduce.Job: Job job_1684298513961_0866
completed successfully
23/05/20 12:10:17 INFO mapreduce.Job: Counters: 54
    File System Counters
        FILE: Number of bytes read=434
        FILE: Number of bytes written=895851
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=474
        HDFS: Number of bytes written=162
        HDFS: Number of read operations=18
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=6
        HDFS: Number of bytes read erasure-coded=0
    Job Counters
        Launched map tasks=1
        Launched reduce tasks=3
        Data-local map tasks=1
```

Total time spent by all maps in occupied slots
(ms)=4465
Total time spent by all reduces in occupied slots
(ms)=24744
Total time spent by all map tasks (ms)=4465
Total time spent by all reduce tasks (ms)=24744
Total vcore-milliseconds taken by all map tasks=4465
Total vcore-milliseconds taken by all reduce
tasks=24744
Total megabyte-milliseconds taken by all map
tasks=4572160
Total megabyte-milliseconds taken by all reduce
tasks=25337856

Map-Reduce Framework

Map input records=13
Map output records=13
Map output bytes=426
Map output materialized bytes=422
Input split bytes=123
Combine input records=0
Combine output records=0
Reduce input groups=6
Reduce shuffle bytes=422
Reduce input records=13
Reduce output records=6
Spilled Records=26
Shuffled Maps =3
Failed Shuffles=0
Merged Map outputs=3
GC time elapsed (ms)=697
CPU time spent (ms)=4550
Physical memory (bytes) snapshot=1158127616
Virtual memory (bytes) snapshot=10370674688
Total committed heap usage (bytes)=1611137024
Peak Map Physical memory (bytes)=526077952
Peak Map Virtual memory (bytes)=2587410432
Peak Reduce Physical memory (bytes)=249159680
Peak Reduce Virtual memory (bytes)=2601119744







Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
Bytes Read=351
File Output Format Counters
Bytes Written=162

creates three folders for three outputs from respective three reducers

Home / user / bigdatalab456422 / training / out9 Trash

<input type="checkbox"/>	Name	Size	User	Group	Permissions	Date
<input type="checkbox"/>			bigdatalab456422	bigdatalab456422	drwxr-xr-x	May 20, 2023 05:09 AM
<input type="checkbox"/>			bigdatalab456422	bigdatalab456422	drwxr-xr-x	May 20, 2023 05:10 AM
<input type="checkbox"/>	 _SUCCESS	0 bytes	bigdatalab456422	bigdatalab456422	-rw-r--r--	May 20, 2023 05:10 AM
<input type="checkbox"/>	 part-r-00000	53 bytes	bigdatalab456422	bigdatalab456422	-rw-r--r--	May 20, 2023 05:10 AM
<input type="checkbox"/>	 part-r-00001	55 bytes	bigdatalab456422	bigdatalab456422	-rw-r--r--	May 20, 2023 05:10 AM
<input type="checkbox"/>	 part-r-00002	54 bytes	bigdatalab456422	bigdatalab456422	-rw-r--r--	May 20, 2023 05:10 AM

Show 45 of 4 itemsPage 1 of 1

for age group0 [age<=20]

Back Home

Page 1 to 1 of 1

Edit file

Refresh

View as binary

Download

/ user / bigdatalab456422 / training / out9 / part-r-00000

1207, Female, bhavya, 20, 15000
1205, Male, kiran, 20, 40000

for age group1 [age > 20 and age <= 30]

Back Home

Page 1 to 1 of 1

Edit file

Refresh

View as binary

/ user / bigdatalab456422 / training / out9 / part-r-00001

1206, Female, laxmi, 25, 35000
1204, Male, prasanth, 30, 31000

for age group2 [age > 30]

[↶ Back](#)

[🏠 Home](#)

Page to of 1 [⏮](#) [⏪](#) [⏩](#) [⏭](#)

[✎ Edit file](#)

[🔄 Refresh](#)

[📄 View as binary](#)

[/ user](#) / [bigdatalab456422](#) / [training](#) / [out9](#) / **part-r-00002**

```
1202, Female, manisha, 40, 51000
1201, Male, gopal, 45, 50000
```