

→ Airlines project in PySpark

```
[bigdatalab456422@ip-10-1-1-204 ~]$ ls -l airlines.csv
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ ls -l airlines.csv
-rw-rw-r-- 1 bigdatalab456422 bigdatalab456422 1821 Jun  5 09:24 airlines.csv
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -put airlines.csv
training
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -put airlines.csv training
[bigdatalab456422@ip-10-1-1-204 ~]$
```

[Back](#) [Home](#) Page 1 to 1 of 1 [⏪](#) [⏩](#) [⏴](#) [⏵](#)

[Edit file](#) [Refresh](#) [View as binary](#) [Download](#)


Last modified06/05/2023 2:55 PM+05:30Userbigdatalab456422Groupbigdatalab456422Size1.78 KBMode100644

/ user / bigdatalab456422 / training / airlines.csv

Year,Quarter,Average revenue per seat,total no. of booked seats
1995,1,296.9,46561
1995,2,296.8,37443
1995,3,287.51,34128
1995,4,287.78,30388
1996,1,283.97,47808
1996,2,275.78,43020
1996,3,269.49,38952
1996,4,278.33,37443
1997,1,283.4,35067
1997,2,289.44,46565
1997,3,282.27,38886
1997,4,293.51,37454
1998,1,304.74,31315
1998,2,300.97,30852
1998,3,315.25,38118
1998,4,316.18,35393
1999,1,331.74,47453
1999,2,329.34,38243

```
[bigdatalab456422@ip-10-1-1-204 ~]$ pyspark
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ pyspark
Python 3.7.6 (default, Jan 8 2020, 19:59:22)
[GC 7.3.0] :: Anaconda, Inc. on linux
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
23/06/05 09:28:38 WARN cluster.YarnSchedulerBackend$YarnSchedulerEndpoint: Attempted to request executors before the AM has registered!
Welcome to
```



version 2.4.0-cdh6.2.1

```
Using Python version 3.7.6 (default, Jan 8 2020 19:59:22)
SparkSession available as 'spark'.
>>>
```

→ find the year with highest revenue

```
>>> from pyspark.sql.types import StructType, StringType,
IntegerType, DoubleType, LongType
```

```
>>> from pyspark.sql.types import StructType, StringType, IntegerType, DoubleType, LongType
>>>
```

```
>>> schema2 =
StructType().add("Year",StringType(),True).add("Quarter",StringType(
,True).add("ARPS",DoubleType(),True).add("Booked_seats",IntegerType(
,True)
```

```
>>> schema2 = StructType().add("Year",StringType(),True).add("Quarter",StringType(),True).add("ARPS",DoubleType(),True).add("Booked_seats",IntegerType(),True)
>>>
```

```
>>> print(schema2)
```

```
>>> print(schema2)
StructType(List(StructField(Year,StringType,true),StructField(Quarter,StringType,true),StructField(ARPS,DoubleType,true),StructField(Booked_seats,IntegerType,true)))
>>>
```

```
>>> df_with_schema2 = spark.read.format("csv").option("header",
"True").schema(schema2).load("hdfs://nameservice1/user/bigdatalab4564
22/training/airlines.csv")
```

```
>>> df_with_schema2 = spark.read.format("csv").option("header", "True").schema(schema2).load("hdfs://nameservice1/user/bigdatalab456422/training/airlines.csv")
>>>
```

```
>>> df_with_schema2.count()
```

```
>>> df_with_schema2.count()
84
>>> █
```

```
>>> df_with_schema2.show()
```

```
>>> df_with_schema2.show()
+-----+
|Year|Quarter|  ARPS|Booked_seats|
+-----+
|1995|    1| 296.9|    46561|
|1995|    2| 296.8|    37443|
|1995|    3| 287.51|    34128|
|1995|    4| 287.78|    30388|
|1996|    1| 283.97|    47808|
|1996|    2| 275.78|    43020|
|1996|    3| 269.49|    38952|
|1996|    4| 278.33|    37443|
|1997|    1| 283.4|    35067|
|1997|    2| 289.44|    46565|
|1997|    3| 282.27|    38886|
|1997|    4| 293.51|    37454|
|1998|    1| 304.74|    31315|
|1998|    2| 300.97|    30852|
|1998|    3| 315.25|    38118|
|1998|    4| 316.18|    35393|
|1999|    1| 331.74|    47453|
|1999|    2| 329.34|    38243|
|1999|    3| 317.22|    33048|
|1999|    4| 317.93|    31256|
+-----+
only showing top 20 rows
>>> █
```

```
>>> df_with_schema2.registerTempTable("airlines")
```

```
>>> df_with_schema2.registerTempTable("airlines")
>>>
```

```
>>> YrWiseRev= spark.sql("SELECT Year, sum(ARPS*Booked_seats) AS
revenue FROM airlines GROUP BY Year ORDER BY revenue DESC")
```

```
>>> YrWiseRev= spark.sql("SELECT Year, sum(ARPS*Booked_seats) AS revenue FROM airlines GROUP BY Year ORDER BY revenue DESC")
>>>
```

```
>>> YrWiseRev.count()
```

```
>>> YrWiseRev.count()
21
>>>
```

```
>>> YrWiseRev.show(21)
```

```
>>> YrWiseRev.show(21)
+-----+
|Year|      revenue|
+-----+
|2013| 6.636320871E7|
|2014| 6.262417585000001E7|
|2015| 6.237899057E7|
|2012| 6.219912728E7|
|2008| 5.7653170760000005E7|
|2007| 5.730921607E7|
|2001| 5.553377999999999E7|
|2010| 5.486152129E7|
|2000| 5.234292655000004E7|
|2011| 5.188828622E7|
|2004| 5.063136494999999E7|
|2006| 5.043789841999999E7|
|2003| 4.927321083E7|
|1999| 4.875771448E7|
|2002| 4.74991465E7|
|2009| 4.674644659E7|
|2005| 4.637678624E7|
|1996| 4.635877803E7|
|1997| 4.538523616E7|
|1995| 4.349424322E7|
|1998| 4.203571778E7|
+-----+

>>>
```

```
>>> YrWiseRev= spark.sql("SELECT Year, sum(ARPS*Booked_seats)/1000000
AS revenue_in_mill FROM airlines GROUP BY Year ORDER BY
revenue_in_mill DESC")
```

```
>>> YrWiseRev= spark.sql("SELECT Year, sum(ARPS*Booked_seats)/1000000 AS revenue_in_mill FROM airlines GROUP BY Year ORDER BY revenue_in_mill DESC")
>>>
```

```
>>> YrWiseRev.count()
```

```
>>> YrWiseRev.count()
21
>>>
```

```
>>> YrWiseRev.show(21)
```

```
>>> YrWiseRev.show(21)
+-----+
|Year|      revenue|
+-----+
|2013| 66.36320871|
|2014| 62.62417585000001|
|2015| 62.37899057|
|2012| 62.19912728|
|2008| 57.65317076|
|2007| 57.30921607|
|2001| 55.53377999999999|
|2010| 54.86152129|
|2000| 52.34292655|
|2011| 51.88828622|
|2004| 50.63136495|
|2006| 50.437898419999999|
|2003| 49.27321083|
|1999| 48.75771448|
|2002| 47.4991465|
|2009| 46.746446590000005|
|2005| 46.37678624|
|1996| 46.35877803|
|1997| 45.38523616|
|1995| 43.49424322|
|1998| 42.03571778|
+-----+

>>>
```

```
>>> YrWiseRev= spark.sql("SELECT Year,
round(sum(ARPS*Booked_seats)/1000000, 2) AS revenue_in_mill FROM
airlines GROUP BY Year ORDER BY revenue_in_mill DESC")
```

```
>>> YrWiseRev= spark.sql("SELECT Year, round(sum(ARPS*Booked_seats)/1000000, 2) AS revenue_in_mill FROM airlines GROUP BY Year ORDER BY revenue_in_mill DESC")
>>>
```

```
>>> YrWiseRev.show(21)
```

```
>>> YrWiseRev.show(21)
+-----+
|Year|revenue_in_mill|
+-----+
|2013|66.36|
|2014|62.62|
|2015|62.38|
|2012|62.2|
|2008|57.65|
|2007|57.51|
|2001|55.53|
|2010|54.86|
|2000|52.34|
|2011|51.89|
|2004|50.63|
|2006|50.44|
|2003|49.27|
|1999|48.76|
|2002|47.5|
|2009|46.75|
|2005|46.38|
|1996|46.36|
|1997|45.39|
|1995|43.49|
|1998|42.04|
+-----+

>>> █
```

→ find the year with highest number of PAX flown

```
>>> YrWisePax = spark.sql("SELECT Year, sum(Booked_seats) AS
total_pax FROM airlines GROUP BY Year ORDER BY total_pax DESC")
```

```
>>> YrWisePax = spark.sql("SELECT Year, sum(Booked_seats) AS total_pax FROM airlines GROUP BY Year ORDER BY total_pax DESC")
>>>
```

```
>>> YrWisePax.count()
```


```
>>> YrWisePax.count()
21
>>>
```

```
>>> YrWisePax.show(21)
```

```
>>> YrWisePax.show(21)
+-----+
|Year|total_pax|
+-----+
|2007|176299|
|2013|173676|
|2001|173598|
|1996|167223|
|2008|166897|
|2012|166076|
|2015|165438|
|2004|164800|
|2010|163741|
|2014|159823|
|1997|157972|
|2003|156153|
|2000|154376|
|2006|153789|
|2002|152195|
|2005|150610|
|2009|150308|
|1999|150000|
|1995|148520|
|2011|142647|
|1998|135678|
+-----+

>>>
```

→ Types of variables in Spark



- a. Accumulators
 - i. Are at driver/spark app
 - ii. Value can change as it sums up the values at each of the worker nodes
 - iii. Active throughout entire session
 - iv. Need to re-initialize it to zero to reset accumulator
- b. Broadcast variables
 - i. Are like constant/immutable variable
 - ii. Can be used in any program, in any stage
 - iii. Declared at driver node, Can be used from any of the worker nodes

→ Note for port number

- a. check core-site.xml
- b. hdfs://localhost:54310

→ Retail Project for accumulator & broadcast variables

→ write a program using shared variables like broadcast and accumulators to calculate total tax and total profit using Retail data

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -ls
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -ls
Found 10 items
drwxr-xr-x - bigdatalab456422 bigdatalab456422 0 2023-05-31 11:00 .Trash
drwxr-xr-x - bigdatalab456422 bigdatalab456422 0 2023-06-05 09:28 .sparkStaging
drwxr-xr-x - bigdatalab456422 bigdatalab456422 0 2023-05-30 11:35 .staging
-rw-r--r-- 3 bigdatalab456422 bigdatalab456422 40990862 2023-06-02 10:33 NYSE.csv
drwxr-xr-x - bigdatalab456422 bigdatalab456422 0 2023-05-19 11:52 data
drwxr-xr-x - bigdatalab456422 bigdatalab456422 0 2023-05-25 07:12 hive
-rw-r--r-- 3 bigdatalab456422 bigdatalab456422 50 2023-05-16 12:43 newfile.txt
drwxr-xr-x - bigdatalab456422 bigdatalab456422 0 2023-05-30 10:34 sales
drwxr-xr-x - bigdatalab456422 bigdatalab456422 0 2023-05-20 10:02 student
drwxr-xr-x - bigdatalab456422 bigdatalab456422 0 2023-06-05 09:25 training

[bigdatalab456422@ip-10-1-1-204 ~]$
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -mkdir retail
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -mkdir retail
[bigdatalab456422@ip-10-1-1-204 ~]$
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -ls retail
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -ls retail
[bigdatalab456422@ip-10-1-1-204 ~]$
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -put D11 D12 D01 D02
retail
```

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop fs -put D11 D12 D01 D02 retail
[bigdatalab456422@ip-10-1-1-204 ~]$
```

```
>>> retailRDD =
sc.textFile("hdfs://nameservice1/user/bigdatalab456422/retail")
>>> retailRDD = sc.textFile("hdfs://nameservice1/user/bigdatalab456422/retail")
>>>
```

```
>>> retailRDD.count()
```

```
>>> retailRDD.count()
817741
>>> █
```

```
>>> retailRDD.getNumPartitions()
```

```
>>> retailRDD.getNumPartitions()
4
```

```
>>> gst = sc.broadcast(5.00)
```

```
>>> gst = sc.broadcast(5.00)
>>>
```

```
>>> totalTax = sc.accumulator(0.00)
```

```
>>> totalTax = sc.accumulator(0.00)
>>>
```

```
>>> totalProfit = sc.accumulator(0.00)
```

```
>>> totalProfit = sc.accumulator(0.00)
>>>
```

```
>>> arrayRDD = retailRDD.map(lambda a : a.split(";"))
```

```
>>> arrayRDD = retailRDD.map(lambda a : a.split(";"))
>>>
```

```
>>> taxAndProfit = arrayRDD.map(lambda a : (
float(a[8])*gst.value/100 , (float(a[8]) - float(a[7]))))
```

```
>>> taxAndProfit = arrayRDD.map(lambda a : ( float(a[8])*gst.value/100 , (float(a[8]) - float(a[7]))))
>>>
```

```
>>> for a in taxAndProfit.take(5):
...     print(a)
>>> for a in taxAndProfit.take(5):
...     print(a)
...
(2.6, 8.0)
(6.45, -21.0)
(1.95, 4.0)
(5.95, 25.0)
(7.95, 59.0)
>>> █
```

```
>>> for line in taxAndProfit.collect():
...     totalTax += line[0]
>>> for line in taxAndProfit.collect():
...     totalTax += line[0]
...
>>>
```

```
>>> print(totalTax)
```

```
>>> print(totalTax)
5392053.600018393
>>> █
```

```
>>> for line in taxAndProfit.collect():
...     totalProfit += line[1]
>>> for line in taxAndProfit.collect():
...     totalProfit += line[1]
...
>>>
```

```
>>> print(totalProfit)
```

```
>>> print(totalProfit)
16163257.0
>>>
```

does not work as totalProfit is of type Accumulator, so it needs to be accessed using value

```
>>> print(totalProfit/1000000)
```

```
>>> print(totalProfit/1000000)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: unsupported operand type(s) for /: 'Accumulator' and 'int'
>>>
```

```
>>> print(totalProfit.value/1000000)
```

```
>>> print(totalProfit.value/1000000)
16.163257
>>>
```

find net profit after deducting tax

```
>>> print(totalProfit.value - totalTax.value)
```

```
>>> print(totalProfit.value - totalTax.value)
10771203.399981607
>>>
```

```
>>> print(totalProfit)
```

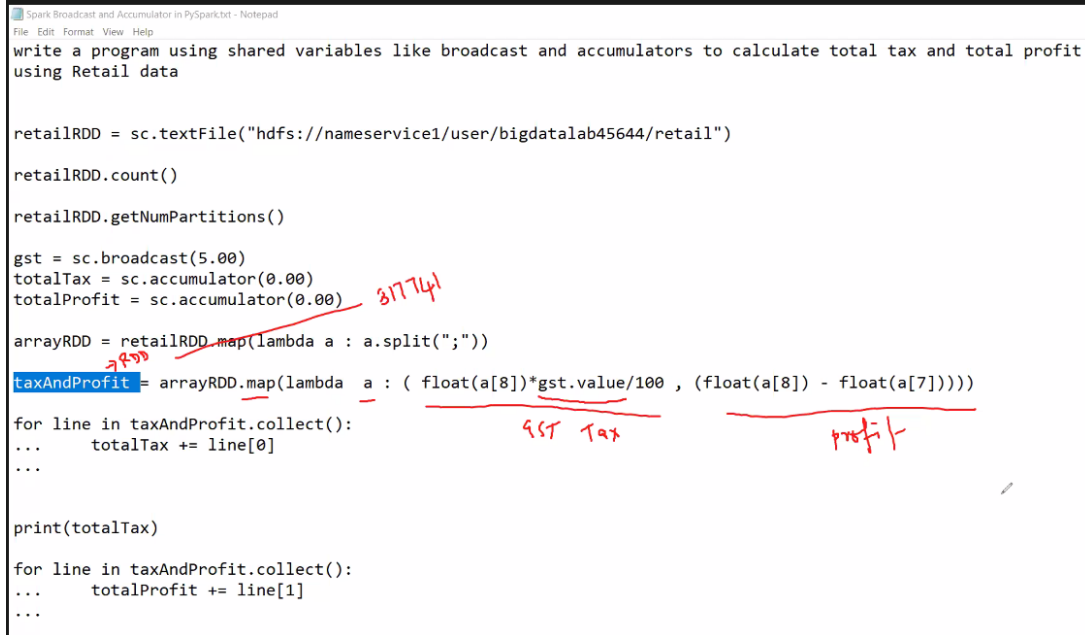
```
>>> print(totalProfit)
16163257.0
>>>
```

```
>>> totalProfit = sc.accumulator(0.0)
```

```
>>> totalProfit = sc.accumulator(0.0)
>>>
```

```
>>> print(totalProfit)
```

```
>>> print(totalProfit)
0.0
>>>
```



```
Spark Broadcast and Accumulator in PySpark.txt - Notepad
File Edit Format View Help

write a program using shared variables like broadcast and accumulators to calculate total tax and total profit
using Retail data

retailRDD = sc.textFile("hdfs://nameservice1/user/bigdatalab45644/retail")

retailRDD.count()

retailRDD.getNumPartitions()

gst = sc.broadcast(5.00)
totalTax = sc.accumulator(0.00)
totalProfit = sc.accumulator(0.00)

arrayRDD = retailRDD.map(lambda a : a.split(";"))

taxAndProfit = arrayRDD.map(lambda a : ( float(a[8])*gst.value/100 , (float(a[8]) - float(a[7]))))

for line in taxAndProfit.collect():
...     totalTax += line[0]
...

print(totalTax)

for line in taxAndProfit.collect():
...     totalProfit += line[1]
...
```

Handwritten red annotations in the image:

- A red line is drawn under the `map` operation in `arrayRDD`.
- A red arrow points to the `taxAndProfit` variable.
- A red line is drawn under the lambda function in `taxAndProfit`.
- Red handwritten text "317741" is written next to the `totalTax` accumulator.
- Red handwritten text "gst Tax" is written under the first part of the lambda function.
- Red handwritten text "profit" is written under the second part of the lambda function.

Previous exam pattern

1. MapReduce Problem 1 no
2. Hive Queries 5 no
3. Spark Queries 5 no