→ YARN cluster - MapReduce Life Cycle

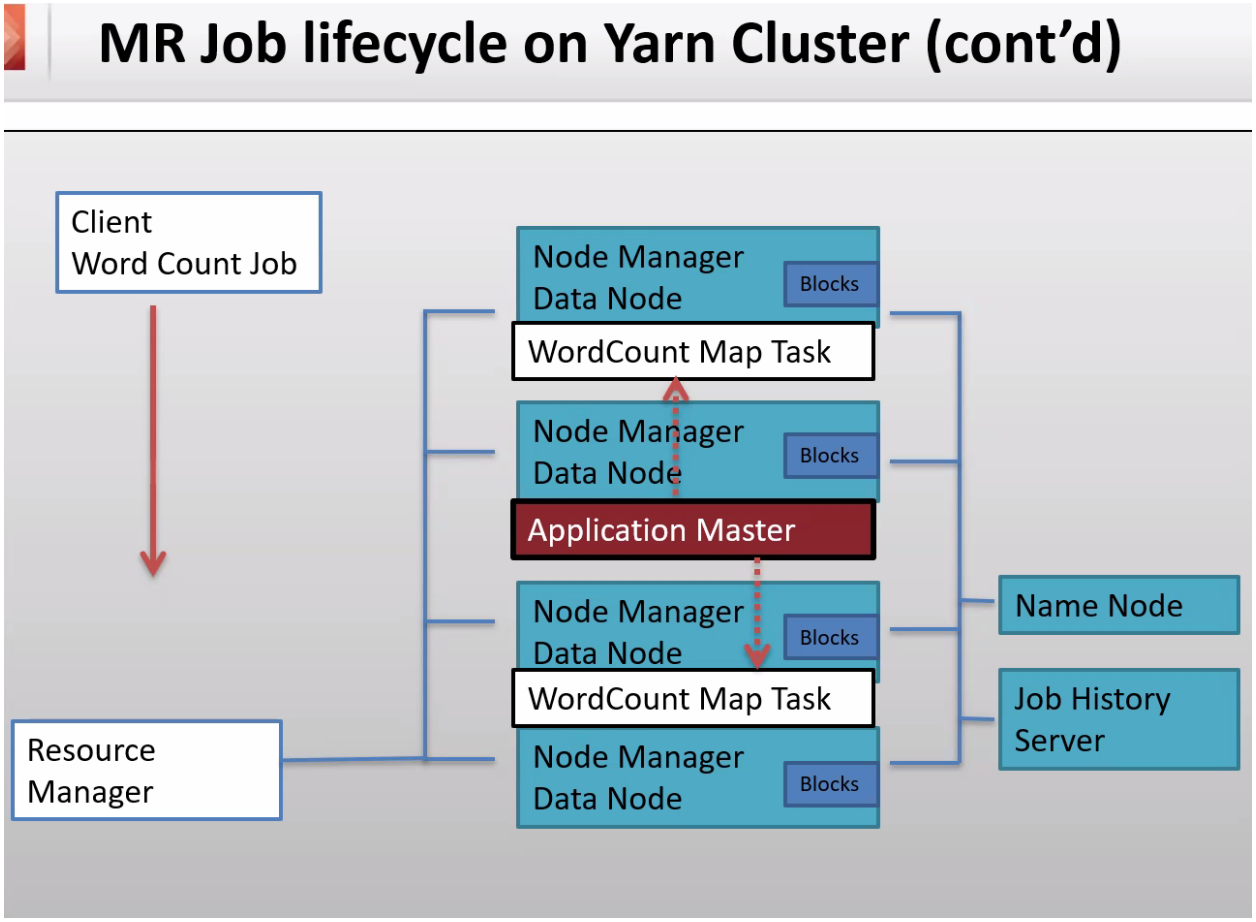

→ Application Master (AM):

a. Is a temporary Service
b. coordinator for jobs
c. Can be on any DataNode containing requested resources
d. Does not process data, but will coordinate the processing, while actual processing is performed by respective Nodemanager on each DataNode
e. Data is processed locally by Nodemanager on each DataNode

→ Map-Reduce

a. Mapper : program to filter the data → key, value
b. Reducer : program to process/aggregate the mapper → key, value

```
Mapper : prog to filter the data -> key, val
Reducer : Prog to process/aggregate the mapper data -> k,v

Retail store example

store id, prod, qty_sold, price/unit,,,,,,


blk_1 [node 1]
1,201,5,80
1,202,10,20
2,201,20,80
2,202,30,20

blk_2 [node 3]
1,201,10,80
1,202,20,20
2,201,30,80
2,202,40,20

prob statement : find total qty sold for each prod

select prod, sum(qty_sold) from sales group by prod;

key = prod
value = qty
```

→ Map-Reduce: Disk Based processing (local disk storage - local node storage)
→ Spark: memory based processing (in-memory storage)

→ Three stages in Map-Reduce:
   a. Map
   b. Shuffle
   c. Reduce

→ Mappers are launched parallel for single job/query
→ Mapper:
   a. Data is filtered
   b. Input is text format, output in key-value format

→ Data shuffling:
   a. First resource container is created for Sort and shuffle stage
   b. Sort & shuffle is processed on separate node
   c. merging and then processing data on another node
   d. Format is key,[list of values from different nodes]
   e. Keys are sorted in ASC order, while values are in random order
   f. Output of shuffling in mapper stage is now input for reducer

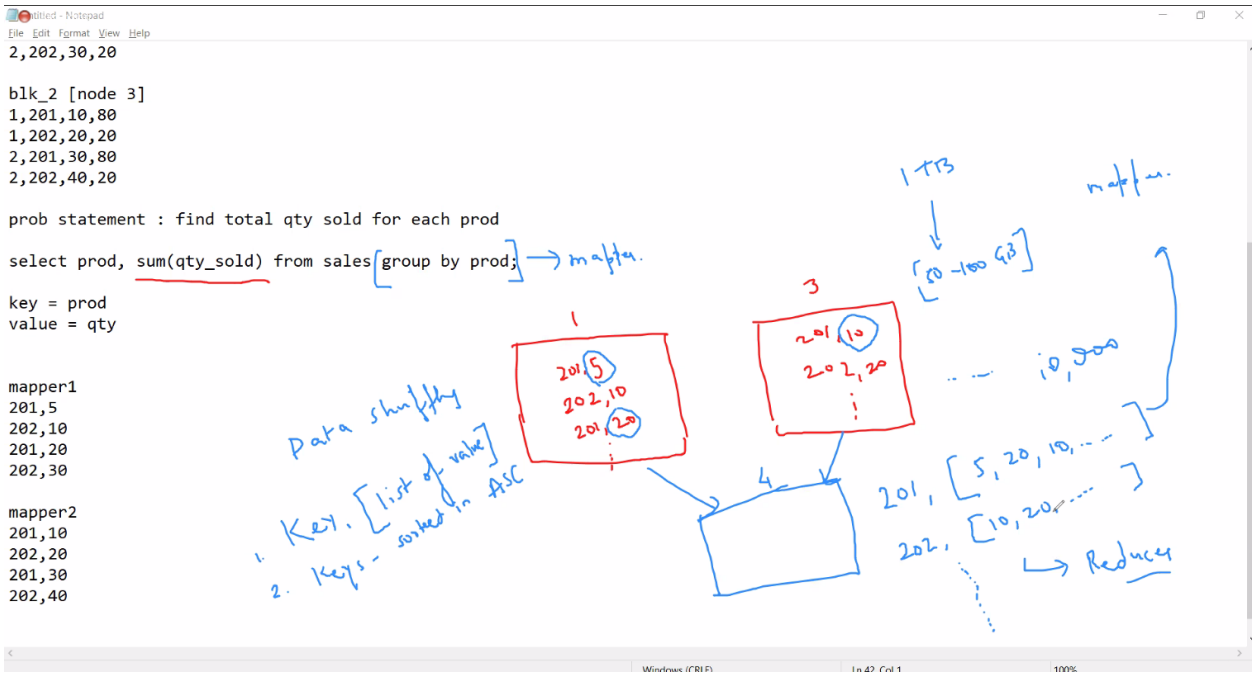

→ Reducer:

a. First resource container is launched for reduce, Then reducer computes the final result And then it writes the final output to the HDFS
b. Reducer is launched only once all the mapping-shuffling is done
c.

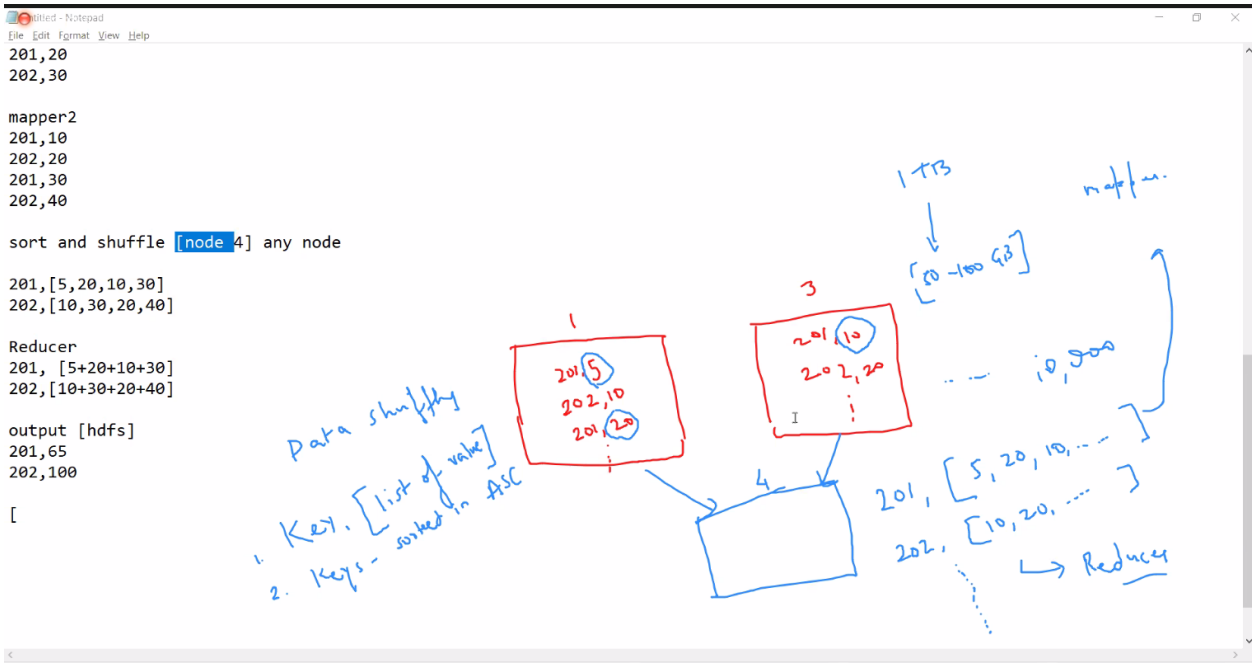

→ find total volume for each stockID

```
Untitled - Notepad
File  Edit  Format  View  Help
202,100

Retail store example

store id, prod, qty_sold, price/unit,,,,,,

blk_1 [node 1]
1,201,5,80
1,202,10,20
2,201,20,80
2,202,30,20

blk_2 [node 3]
1,201,10,80
1,202,20,20
2,201,30,80
2,202,40,20

prob statement : find total sales for each store id

select store_id, sum(qty*price) from sales group by store_id

key = store_id
val = price*qty

m1
```

```
Untitled - Notepad
File  Edit  Format  View  Help
key = store_id
val = price*qty

m1
1,400
1,200
2,1600
2,600

m2
1,800
1,400
2,2400
2,800

1, [400,200,800,400]
2, [1600,600,2400,800]

Reducer
--------
1, [400+200+800+400]
2, [1600+600+2400+800]


output
1,1800
2,5400
```

first upload myjar.jar file using ftp

```
[bigdatalab456422@ip-10-1-1-204 ~]$ ll
total 959100
-rw-rw-r-- 1 bigdatalab456422 bigdatalab456422 207106008 May 17 12:44
eclipse.gz
```

```
-rw-rw-r-- 1 bigdatalab456422 bigdatalab456422        50 May 16 12:19
file1.txt
-rw-rw-r-- 1 bigdatalab456422 bigdatalab456422        20 May 16 12:30
file2.txt
-rw-rw-r-- 1 bigdatalab456422 bigdatalab456422 209715200 May 17 09:16
myfile
-rw-rw-r-- 1 bigdatalab456422 bigdatalab456422 524288000 May 17 09:35
myfile2
-rw-rw-r-- 1 bigdatalab456422 bigdatalab456422      4088 May 18 11:41
myjar.jar
-rw-rw-r-- 1 bigdatalab456422 bigdatalab456422  40990862 May 17 09:21
NYSE.csv

[bigdatalab456422@ip-10-1-1-204 ~]$ jar tvf myjar.jar
   25 Thu May 18 17:09:20 UTC 2023 META-INF/MANIFEST.MF
  387 Thu May 18 15:53:20 UTC 2023 .project
 2408 Thu May 18 17:00:02 UTC 2023 StockVolume$MapClass.class
 2349 Thu May 18 17:00:02 UTC 2023 StockVolume$ReduceClass.class
 1697 Thu May 18 17:00:02 UTC 2023 StockVolume.class
  640 Thu May 18 17:00:00 UTC 2023 .classpath
```

run map & reduce task
generates part-r-0000 file, ‘r’ means reducer output

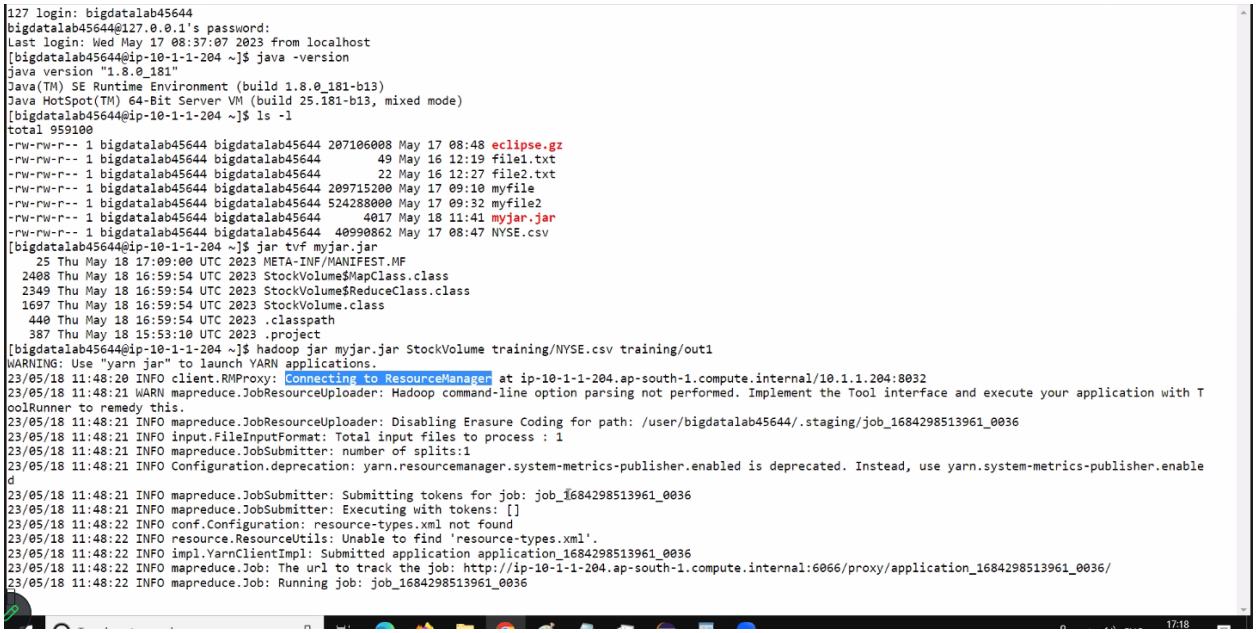

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop jar myjar.jar StockVolume
training/NYSE.csv training/out1
WARNING: Use "yarn jar" to launch YARN applications.
23/05/18 11:48:33 INFO client.RMProxy: Connecting to ResourceManager
at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
23/05/18 11:48:33 WARN mapreduce.JobResourceUploader: Hadoop
command-line option parsing not performed. Implement the Tool
interface and execute your application with T
oolRunner to remedy this.
23/05/18 11:48:33 INFO mapreduce.JobResourceUploader: Disabling
Erasure Coding for path:
/user/bigdatalab456422/.staging/job_1684298513961_0041
23/05/18 11:48:34 INFO input.FileInputFormat: Total input files to
process : 1
23/05/18 11:48:34 INFO mapreduce.JobSubmitter: number of splits:1
23/05/18 11:48:34 INFO Configuration.deprecation:
yarn.resourcemanager.system-metrics-publisher.enabled is deprecated.
Instead, use yarn.system-metrics-publisher.enable
d
23/05/18 11:48:34 INFO mapreduce.JobSubmitter: Submitting tokens for
job: job_1684298513961_0041
23/05/18 11:48:34 INFO mapreduce.JobSubmitter: Executing with tokens:
[]
23/05/18 11:48:34 INFO conf.Configuration: resource-types.xml not
found
23/05/18 11:48:34 INFO resource.ResourceUtils: Unable to find
'resource-types.xml'.
23/05/18 11:48:34 INFO impl.YarnClientImpl: Submitted application
application_1684298513961_0041
23/05/18 11:48:34 INFO mapreduce.Job: The url to track the job:
http://ip-10-1-1-204.ap-south-1.compute.internal:6066/proxy/applicati
on_1684298513961_0041/
23/05/18 11:48:34 INFO mapreduce.Job: Running job:
job_1684298513961_0041
23/05/18 11:48:58 INFO mapreduce.Job: Job job_1684298513961_0041
running in uber mode : false
23/05/18 11:48:58 INFO mapreduce.Job:  map 0% reduce 0%
23/05/18 11:49:27 INFO mapreduce.Job:  map 67% reduce 0%
23/05/18 11:49:29 INFO mapreduce.Job:  map 100% reduce 0%
23/05/18 11:50:00 INFO mapreduce.Job:  map 100% reduce 100%
23/05/18 11:50:01 INFO mapreduce.Job: Job job_1684298513961_0041
completed successfully
23/05/18 11:50:01 INFO mapreduce.Job: Counters: 54
        File System Counters
```

```
                FILE: Number of bytes read=3584395
                FILE: Number of bytes written=7613961
                FILE: Number of read operations=0
                FILE: Number of large read operations=0
                FILE: Number of write operations=0
                HDFS: Number of bytes read=40990986
                HDFS: Number of bytes written=2918
                HDFS: Number of read operations=8
                HDFS: Number of large read operations=0
                HDFS: Number of write operations=2
                HDFS: Number of bytes read erasure-coded=0
        Job Counters
                Launched map tasks=1
                Launched reduce tasks=1
                Data-local map tasks=1
                Total time spent by all maps in occupied slots
(ms)=26310
                Total time spent by all reduces in occupied slots
(ms)=28739
                Total time spent by all map tasks (ms)=26310
                Total time spent by all reduce tasks (ms)=28739
                Total vcore-milliseconds taken by all map tasks=26310
                Total vcore-milliseconds taken by all reduce
tasks=28739
                Total megabyte-milliseconds taken by all map
tasks=26941440
                Total megabyte-milliseconds taken by all reduce
tasks=29428736
        Map-Reduce Framework
                Map input records=735026
                Map output records=735026
                Map output bytes=8781587
                Map output materialized bytes=3584391
                Input split bytes=124
                Combine input records=0
                Combine output records=0
                Reduce input groups=203
                Reduce shuffle bytes=3584391
                Reduce input records=735026
                Reduce output records=203
                Spilled Records=1470052
                Shuffled Maps =1
                Failed Shuffles=0
                Merged Map outputs=1
                GC time elapsed (ms)=575
```

```
        CPU time spent (ms)=8990
        Physical memory (bytes) snapshot=882450432
        Virtual memory (bytes) snapshot=5186535424
        Total committed heap usage (bytes)=1075314688
        Peak Map Physical memory (bytes)=621236224
        Peak Map Virtual memory (bytes)=2586968064
        Peak Reduce Physical memory (bytes)=261267456
        Peak Reduce Virtual memory (bytes)=2599567360
    Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0
    File Input Format Counters
        Bytes Read=40990862
    File Output Format Counters
        Bytes Written=2918
```

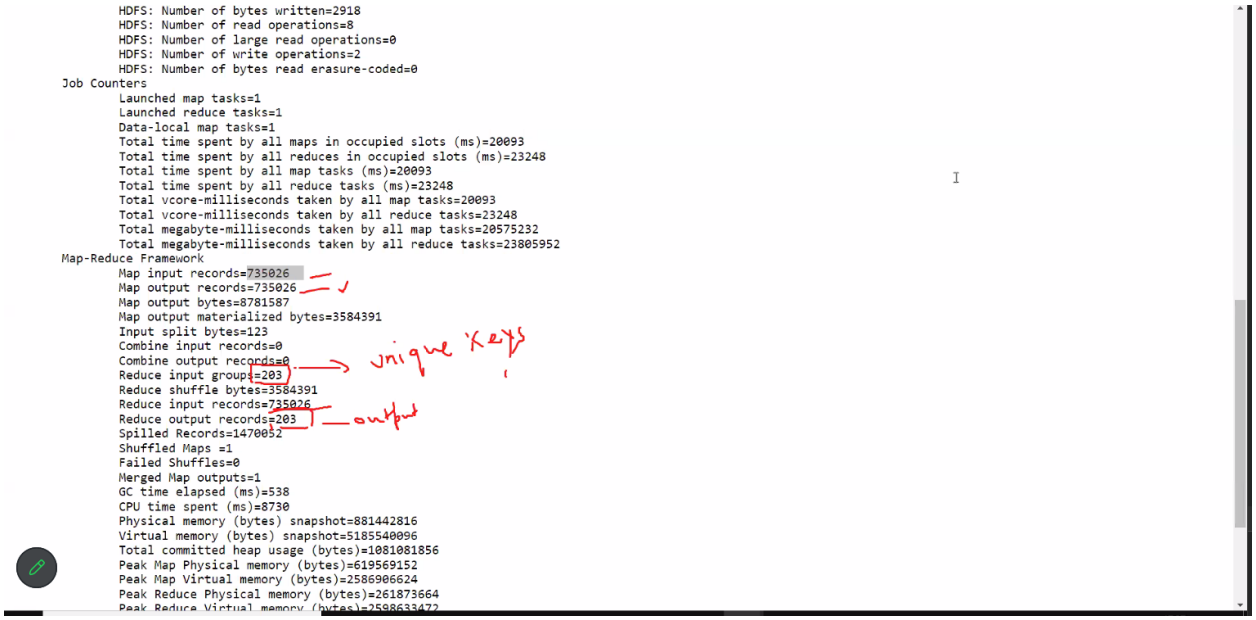

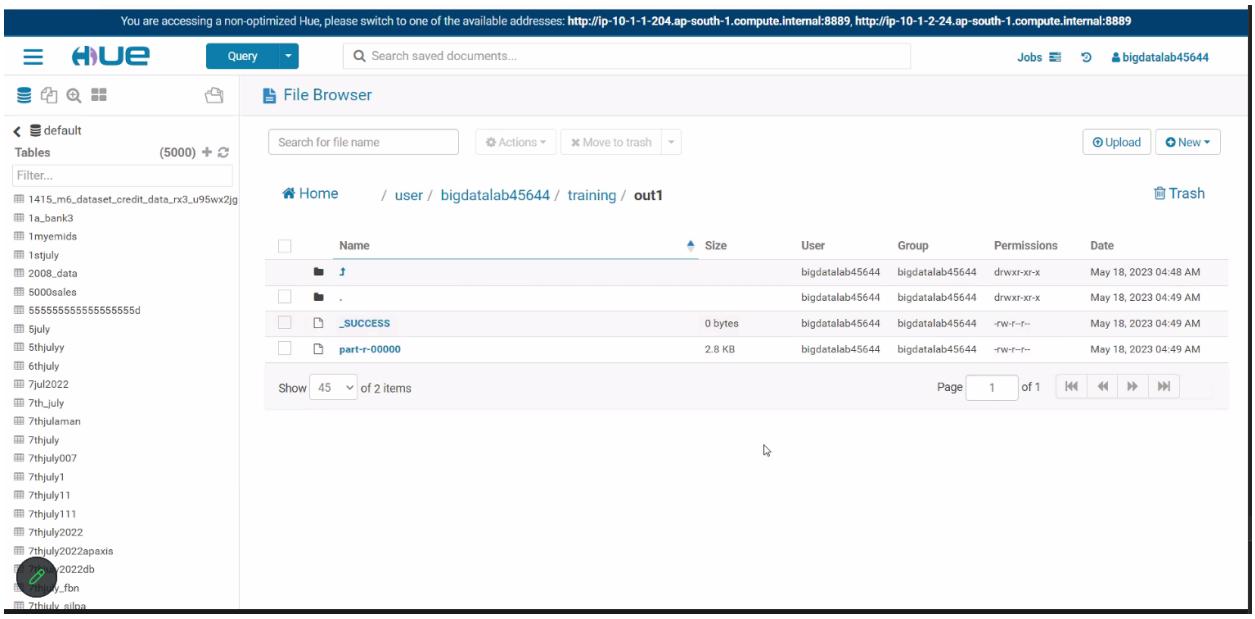
You are accessing a non-optimized Hue, please switch to one of the available addresses: http://ip-10-1-1-204.ap-south-1.compute.internal:8889, http://ip-10-1-2-24.ap-south-1.compute.internal:8889
Query
Search saved documents...
Jobs
bigdatalab45644
File Browser
default
Tables (5000)
Filter...
1415_m6_dataset_credit_data_rx3_u95wx2jg
1a_bank3
1myemids
1stjuly
2008_data
5000sales
55555555555555555d
5july
5thjulyy
6thjuly
7jul2022
7th_july
7thjulaman
7thjuly
7thjuly007
7thjuly1
7thjuly11
7thjuly111
7thjuly2022
7thjuly2022apaxis
7thjuly2022db
7thjuly_fbn
7thjuly_silpa
Search for file name
Actions
Move to trash
Upload
New
Home / user / bigdatalab45644 / training / out1
Trash
Name
Size
User
Group
Permissions
Date
bigdatalab45644 bigdatalab45644 drwxr-xr-x May 18, 2023 04:48 AM
. bigdatalab45644 bigdatalab45644 drwxr-xr-x May 18, 2023 04:49 AM
_SUCCESS 0 bytes bigdatalab45644 bigdatalab45644 -rw-r--r-- May 18, 2023 04:49 AM
part-r-00000 2.8 KB bigdatalab45644 bigdatalab45644 -rw-r--r-- May 18, 2023 04:49 AM
Show 45 of 2 items
Page 1 of 1

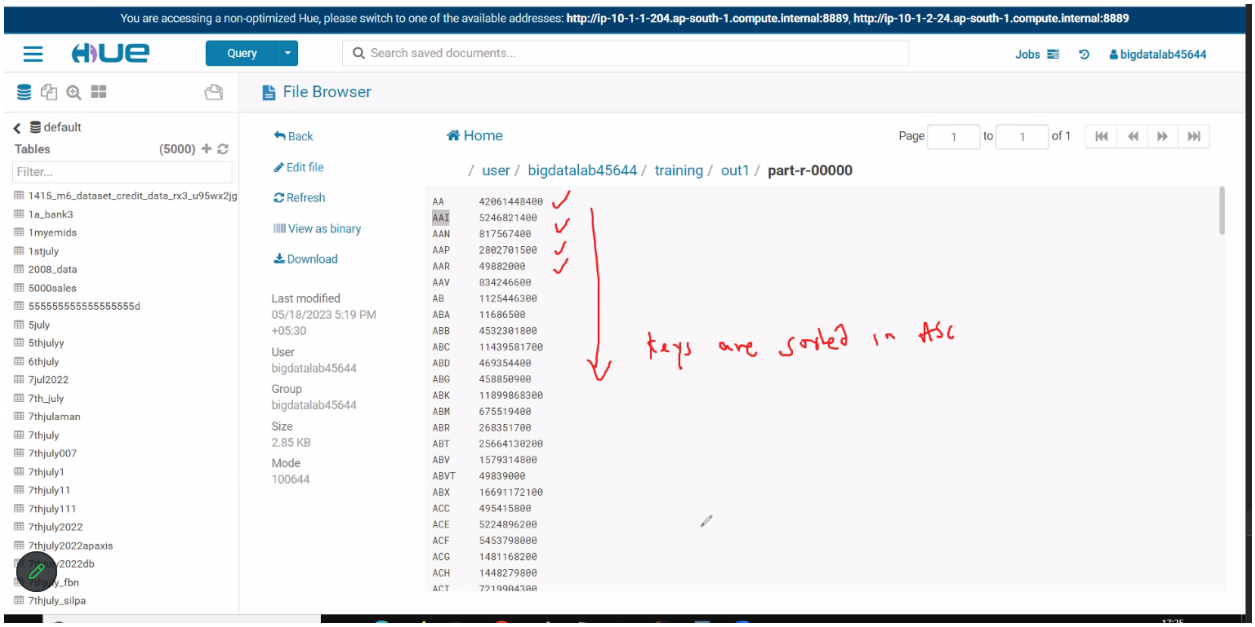
You are accessing a non-optimized Hue, please switch to one of the available addresses: http://ip-10-1-1-204.ap-south-1.compute.internal:8889, http://ip-10-1-2-24.ap-south-1.compute.internal:8889
Query
Search saved documents...
Jobs
bigdatalab45644
File Browser
default
Tables (5000)
Filter...
1415_m6_dataset_credit_data_rx3_u95wx2jg
1a_bank3
1myemids
1stjuly
2008_data
5000sales
55555555555555555d
5july
5thjulyy
6thjuly
7jul2022
7th_july
7thjulaman
7thjuly
7thjuly007
7thjuly1
7thjuly11
7thjuly111
7thjuly2022
7thjuly2022apaxis
7thjuly2022db
7thjuly_fbn
7thjuly_silpa
Back
Edit file
Refresh
View as binary
Download
Last modified
05/18/2023 5:19 PM
+05:30
User
bigdatalab45644
Group
bigdatalab45644
Size
2.85 KB
Mode
100644
Home
Page 1 to 1 of 1
/ user / bigdatalab45644 / training / out1 / part-r-00000
AA 42061448400
AAI 5246821400
AAN 817567400
AAP 2802701500
AAR 49882000
AAV 834246600
AB 1125446300
ABA 11686500
ABB 4532301800
ABC 11439581700
ABD 469354400
ABG 458850900
ABK 11899868300
ABM 675519400
ABR 268351700
ABT 25664130200
ABV 1579314800
ABVT 49839000
ABX 16691172100
ACC 495415800
ACE 5224896200
ACF 5453798000
ACG 1481168200
ACH 1448279800
keys are sorted in asc

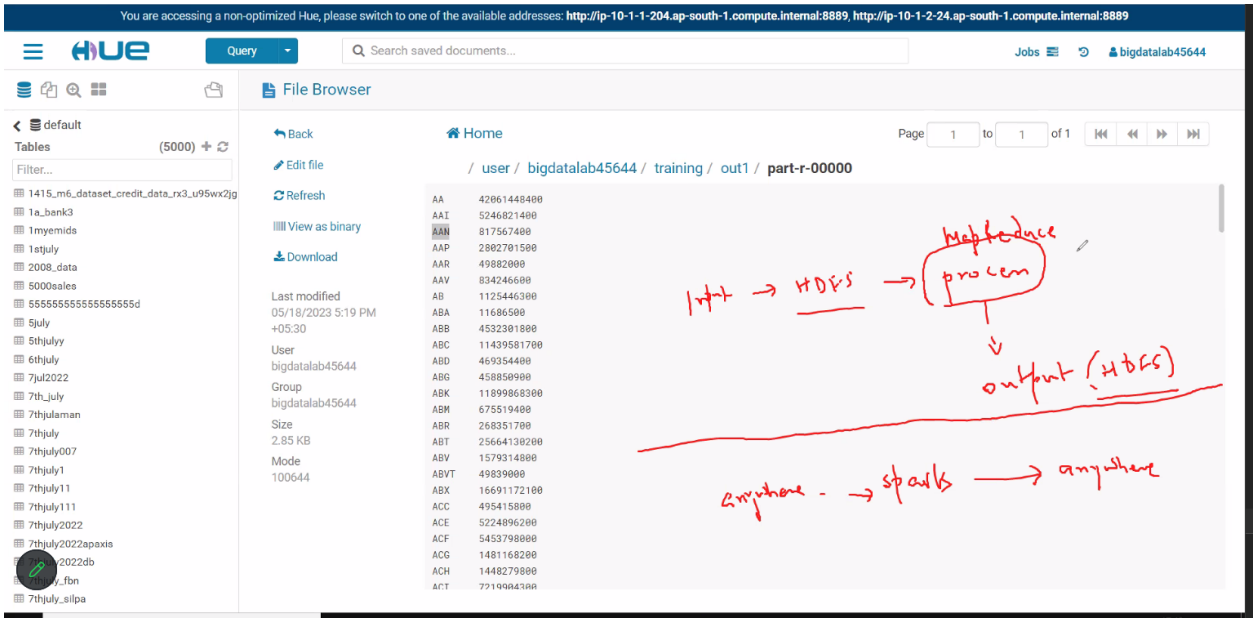

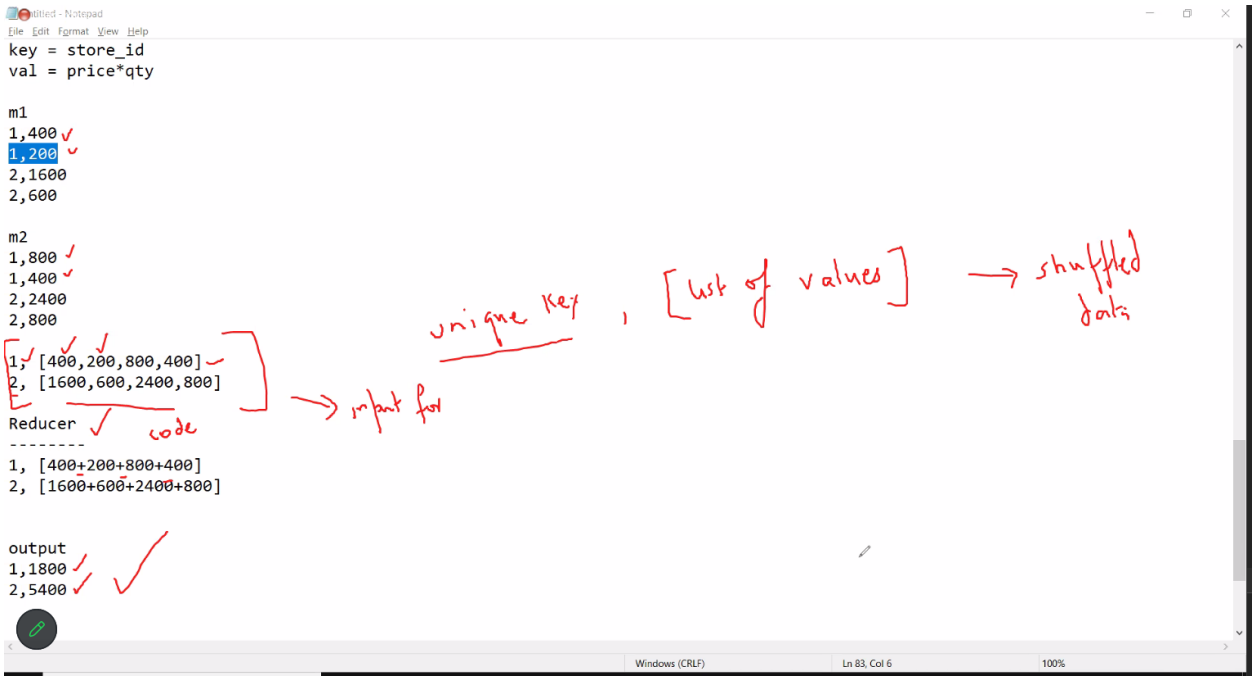

→

run only map task

change java code for reducerTask as '0'

```
job.setNumReduceTasks(0);
```

generates part-r-0000 file, 'r' means reducer output

```
[bigdatalab456422@ip-10-1-1-204 ~]$ hadoop jar myjar.jar StockVolume
training/NYSE.csv training/out2
WARNING: Use "yarn jar" to launch YARN applications.
23/05/18 12:23:09 INFO client.RMProxy: Connecting to ResourceManager
at ip-10-1-1-204.ap-south-1.compute.internal/10.1.1.204:8032
23/05/18 12:23:10 WARN mapreduce.JobResourceUploader: Hadoop
command-line option parsing not performed. Implement the Tool
interface and execute your application with T
oolRunner to remedy this.
23/05/18 12:23:10 INFO mapreduce.JobResourceUploader: Disabling
Erasure Coding for path:
/user/bigdatalab456422/.staging/job_1684298513961_0137
23/05/18 12:23:10 INFO input.FileInputFormat: Total input files to
process : 1
23/05/18 12:23:10 INFO mapreduce.JobSubmitter: number of splits:1
23/05/18 12:23:10 INFO Configuration.deprecation:
yarn.resourcemanager.system-metrics-publisher.enabled is deprecated.
Instead, use yarn.system-metrics-publisher.enable
d
23/05/18 12:23:10 INFO mapreduce.JobSubmitter: Submitting tokens for
job: job_1684298513961_0137
23/05/18 12:23:10 INFO mapreduce.JobSubmitter: Executing with tokens:
[]
23/05/18 12:23:11 INFO conf.Configuration: resource-types.xml not
found
23/05/18 12:23:11 INFO resource.ResourceUtils: Unable to find
'resource-types.xml'.
23/05/18 12:23:11 INFO impl.YarnClientImpl: Submitted application
application_1684298513961_0137
23/05/18 12:23:11 INFO mapreduce.Job: The url to track the job:
http://ip-10-1-1-204.ap-south-1.compute.internal:6066/proxy/applicati
on_1684298513961_0137/
23/05/18 12:23:11 INFO mapreduce.Job: Running job:
job_1684298513961_0137
23/05/18 12:23:20 INFO mapreduce.Job: Job job_1684298513961_0137
running in uber mode : false
23/05/18 12:23:20 INFO mapreduce.Job:  map 0% reduce 0%
23/05/18 12:23:26 INFO mapreduce.Job:  map 100% reduce 0%
23/05/18 12:23:26 INFO mapreduce.Job: Job job_1684298513961_0137
completed successfully
23/05/18 12:23:26 INFO mapreduce.Job: Counters: 33
       File System Counters
               FILE: Number of bytes read=0
               FILE: Number of bytes written=222430
               FILE: Number of read operations=0
```

```
                FILE: Number of large read operations=0
                FILE: Number of write operations=0
                HDFS: Number of bytes read=40990986
                HDFS: Number of bytes written=7842509
                HDFS: Number of read operations=7
                HDFS: Number of large read operations=0
                HDFS: Number of write operations=2
                HDFS: Number of bytes read erasure-coded=0
        Job Counters
                Launched map tasks=1
                Data-local map tasks=1
                Total time spent by all maps in occupied slots
(ms)=4102
                Total time spent by all reduces in occupied slots
(ms)=0
                Total time spent by all map tasks (ms)=4102
                Total vcore-milliseconds taken by all map tasks=4102
                Total megabyte-milliseconds taken by all map
tasks=4200448
        Map-Reduce Framework
                Map input records=735026
                Map output records=735026
                Input split bytes=124
                Spilled Records=0
                Failed Shuffles=0
                Merged Map outputs=0
                GC time elapsed (ms)=69
                CPU time spent (ms)=3170
                Physical memory (bytes) snapshot=361504768
                Virtual memory (bytes) snapshot=2584989696
                Total committed heap usage (bytes)=480772096
                Peak Map Physical memory (bytes)=361504768
                Peak Map Virtual memory (bytes)=2584989696
        File Input Format Counters
                Bytes Read=40990862
        File Output Format Counters
                Bytes Written=7842509
```

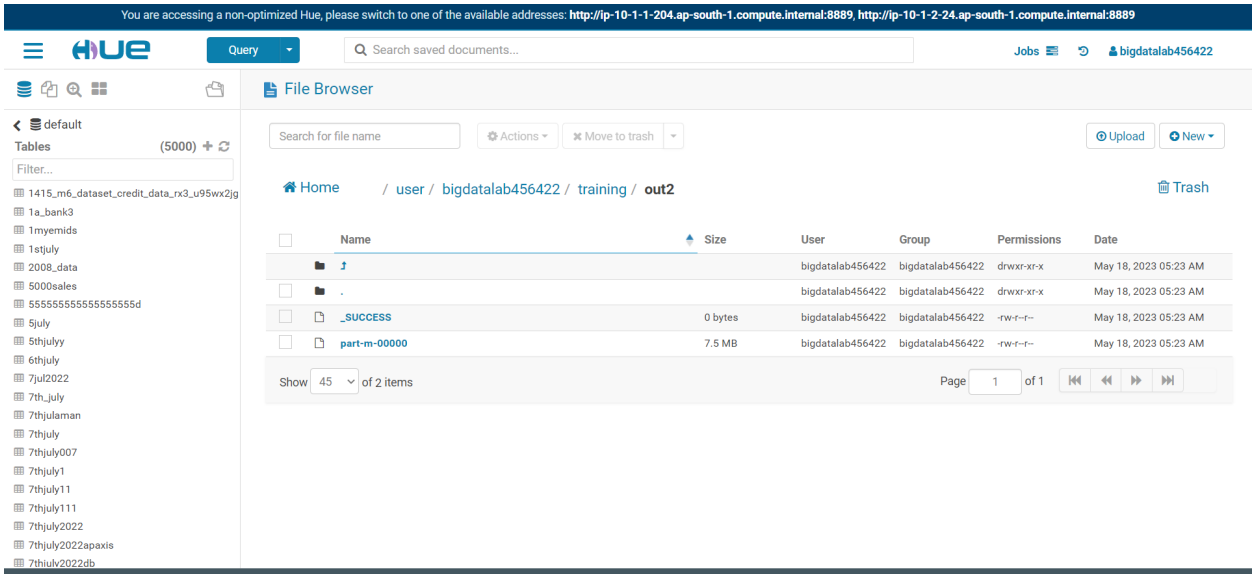

here , in part-m-0000, ‘m’ means output from mapper