

▼ Logistic Regression Model

1. Binary Logistic Regression Model
2. Nominal Logistic Regression Model
3. Ordinal Logistic Regression Model

Binary Logistic Regression Model

- Binary Logistic equation
- when the response has binary values

▼ Binary Logistic Regression Model (mtcars dataset)

▼ import statsmodels

```
1  import numpy as np
2  import pandas as pd
3  # import numpy & pandas
4
5  import statsmodels
6  import statsmodels.api as sm
7  from statsmodels.formula.api import ols
8  import statsmodels.stats.multicomp
9  # import statsmodels
10
11 import sklearn
12 from sklearn.model_selection import train_test_split
13 from sklearn.metrics import confusion_matrix
14 from sklearn.metrics import classification_report
15 # import sklearn
```

▼ upload dataset

```
1 from google.colab import files
2 uploaded=files.upload()
3 # CDAC_DataBook.xlsx
4 # to be used with google colab
5
6 # import os
```

```

7 # os.chdir(r'C:\Users\surya\Downloads\PG-DBDA-Mar23\Datasets')
8 # os.getcwd()
9 # to change current working directory to specified path
10 # to be used while running on local system

```

[Choose Files](#) No file chosen

Upload widget is only available when the cell has been executed in

the current browser session. Please rerun this cell to enable.

Saving CDAC DataBook.xlsx to CDAC DataBook (1).xlsx

▼ `pd.read_excel('WorkBook_Name.xlsx', sheet_name='SheetName')`

```

1 df = pd.read_excel('CDAC_DataBook.xlsx', sheet_name='mtcars')
2 # reading excel file with specifying the sheet name to load dataset
3 df.head()

```

	mpg	cyl	dis	hp	drat	wt	qsec	vs	am	gear	carb
0	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

```

1 df = df[['hp', 'wt', 'am']]
2 # selecting specific columns from dataset
3 df.head()
4 # printing new DataFrame of only required columns

```

	hp	wt	am
0	110	2.620	1
1	110	2.875	1
2	93	2.320	1
3	110	3.215	0
4	175	3.440	0

▼ `df.drop('Pred_col',axis=1)`

- drops response col

```

1 x_train = df.drop('am',axis=1)
2 # manually creating Predictor-Training DataFrame by dropping Response col
3 # without train_test_split()
4 # x_test is not created
5 x_train.head()

```

	hp	wt
0	110	2.620
1	110	2.875
2	93	2.320
3	110	3.215
4	175	3.440

▼ df['Response']

- selects response col

```

1 y_train = df['am']
2 # manually creating Response-Training Series by selecting only Response col
3 # without train_test_split()
4 # y_test is not created
5 y_train.head()

```

```

0    1
1    1
2    1
3    0
4    0
Name: am, dtype: int64

```

▼ sm.add_constant(x_train,prepend=False)

```

1 x_train=sm.add_constant(x_train,prepend=False)
2 # adding constant to match equation
3 # if we do not add constant, model will follow Binary Logistic Equation but without a constant
4 # if we add constant, model will follow Binary Logistic Equation with constant
5 x_train.head()
6 # checking columns in Predictor-training DataFrame after adding constant

```

	hp	wt	const
0	110	2.620	1.0
1	110	2.875	1.0
2	93	2.320	1.0
3	110	3.215	1.0
4	175	3.440	1.0

▼ sm.Logit(y_train,x_train).fit()

```
1 mod1=sm.Logit(y_train,x_train).fit()
2 # creating model using sm.Logit(y, x).fit() method
```

```
Optimization terminated successfully.
      Current function value: 0.157174
      Iterations 9
```

```
1 print(mod1.summary())
2 # printing model summary / Logit Regression Results
```

```

                        Logit Regression Results
=====
Dep. Variable:          am      No. Observations:          32
Model:                  Logit      Df Residuals:            29
Method:                  MLE        Df Model:              2
Date:                   Fri, 30 Jun 2023      Pseudo R-squ.:        0.7673
Time:                   09:47:30      Log-Likelihood:       -5.0296
converged:              True      LL-Null:             -21.615
Covariance Type:        nonrobust      LLR p-value:         6.267e-08
=====
               coef      std err          z      P>|z|      [0.025      0.975]
-----
hp              0.0363      0.018      2.044      0.041      0.001      0.071
wt             -8.0835      3.069     -2.634      0.008     -14.098     -2.069
const          18.8663      7.444      2.535      0.011       4.277     33.455
=====
```

Possibly complete quasi-separation: A fraction 0.12 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

▼ pd.DataFrame([list1, list2, list3], columns=index)

```

1 mydata = pd.DataFrame([[120, 2, 1], [120, 2, 5], [120, 3, 1]], columns=['hp', 'wt', 'const'])
2 # manually creating a Predictor-Test DataFrame using random data
3
4 # DataFrame created keeping hp constant but varying 'wt' col to
5 # see influence of wt on 'am' col
6
7 mydata
8 # printing Predictor-Test DataFrame

```

	hp	wt	const
0	120	2	1
1	120	2	5
2	120	3	1

▼ model.predict(Test_DataFrame)

```

1 mod1.predict(mydata)
2 # generating Prediction Series

```

```

0    0.999133
1    1.000000
2    0.262415
dtype: float64

```

▼ $\text{np.log}(\text{Pred}_a / (1 - \text{Pred}_a)) - \text{np.log}(\text{Pred}_b / (1 - \text{Pred}_b))$

- Binary Logistic Regression Equation

```

1 np.log(0.999133/(1-0.999133)) - np.log(0.262415/(1-0.262415))
2 # np.log(Pred_a/(1-Pred_a)) - np.log(Pred_b/(1-Pred_b))
3 # follows Binary Logistic Regression Equation to generate coefficient
4
5 # Pred_a : take prediction for the lesser 'wt' of Test_DataFrame
6 # Pred_b : take prediction for the highest 'wt' of Test_DataFrame
7
8
9 # find difference of log of odds of two predictions
10 # odds = p(1-P)

```

```
8.083058320861667
```

- Binary Logistic Regression Equation/coefficient = difference of logs of odds of two predictors
- Binary Logistic Regression Equation/coefficient = $\log(\text{odds of A}) - \log(\text{odds of B})$
- Binary Logistic Regression Equation/coefficient = $\text{np.log}(\text{Pred_a}/(1-\text{Pred_a})) - \text{np.log}(\text{Pred_b}/(1-\text{Pred_b}))$
- $\text{odds} = p/(1-p)$
- $\text{odds} = \text{Proportion}(1-\text{Proportion})$
- Pred_a : take prediction for the lesser 'wt' of Test_DataFrame
- Pred_b : take prediction for the highest 'wt' of Test_DataFrame

▼ Binary Logistic Regression Model (diabetes dataset)

▼ import statsmodels & sklearn

```

1 import numpy as np
2 import pandas as pd
3 # import numpy & pandas
4
5 import statsmodels
6 import statsmodels.api as sm
7 from statsmodels.formula.api import ols
8 import statsmodels.stats.multicomp
9 # import statsmodels
10
11 import sklearn
12 from sklearn.model_selection import train_test_split
13 from sklearn.metrics import confusion_matrix
14 from sklearn.metrics import classification_report
15 # import sklearn

```

▼ upload dataset

```

1 from google.colab import files
2 uploaded=files.upload()
3 # CDAC_DataBook.xlsx
4 # to be used with google colab
5
6 # import os
7 # os.chdir(r'C:\Users\surya\Downloads\PG-DBDA-Mar23\Datasets')

```

```

8 # os.getcwd()
9 # to change current working directory to specified path
10 # to be used while running on local system

```

Choose Files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
 Saving CDAC DataBook.xlsx to CDAC DataBook.xlsx

▼ `pd.read_excel('WorkBook_Name.xlsx', sheet_name='SheetName')`

```

1 df = pd.read_excel('CDAC_DataBook.xlsx', sheet_name='diabetes')
2 # reading excelfile with specifying the sheet name to load dataset
3 df.head()

```

	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Out
0	148	72	35	0	33.6	0.63	50	
1	85	66	29	0	26.6	0.35	31	
2	183	64	0	0	23.3	0.67	32	
3	89	66	23	94	28.1	0.17	21	
4	137	40	35	168	43.1	2.29	33	

▼ `df.columns`

```

1 df.columns
2 # fetching column names

```

```

Index(['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI',
      'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')

```

```

1 df = df[['Glucose', 'BloodPressure', 'Age', 'Outcome']]
2 # selecting specific columns from dataset
3 df.head()
4 # printing new DataFrame of only required columns

```

	Glucose	BloodPressure	Age	Outcome
0	148	72	50	1
1	85	66	31	0

▼ `train_test_split(predictor_cols, response_col, test_size=0.25)`

```
1 x_train, x_test, y_train, y_test = train_test_split(df.drop('Outcome', axis=1), df['Outcome'], test_size=.025)
2 # splitting dataset 4-ways
3 # rows are randomly selected for testing
```

▼ `sm.add_constant(x_train, prepend=False)`

```
1 x_train = sm.add_constant(x_train, prepend=False)
2 # adding constant to match equation
3 # if we do not add constant, model will follow Binary Logistic Equation but without a constant
4 # if we add constant, model will follow Binary Logistic Equation with constant
```

▼ `sm.OLS(y_train, x_train).fit()`

```
1 mod1 = sm.OLS(y_train, x_train).fit()
2 # creating model using sm.OLS().fit()
```

```
1 print(mod1.summary())
2 # printing model summary / OLS Regression Results
```

```

OLS Regression Results
=====
Dep. Variable:      Outcome      R-squared:      0.235
Model:              OLS          Adj. R-squared: 0.232
Method:             Least Squares  F-statistic:   76.17
Date:               Thu, 29 Jun 2023  Prob (F-statistic): 5.65e-43
Time:               19:24:03      Log-Likelihood: -408.23
No. Observations:   748          AIC:            824.5
Df Residuals:       744          BIC:            842.9
Df Model:           3
Covariance Type:    nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Glucose	0.0066	0.000	13.221	0.000	0.006	0.008
BloodPressure	-0.0008	0.001	-0.942	0.346	-0.002	0.001
Age	0.0051	0.001	3.693	0.000	0.002	0.008


```

const          -0.5607      0.079      -7.118      0.000      -0.715      -0.406
=====
Omnibus:                51.943      Durbin-Watson:                1.904
Prob(Omnibus):           0.000      Jarque-Bera (JB):           42.118
Skew:                   0.495      Prob(JB):                 7.15e-10
Kurtosis:               2.391      Cond. No.                 756.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Establishing Null Hypothesis H_0

- H_0 : BloodPressure does not impact outcome
- BloodPresure P-Value 0.346 > 0.05, so we do not reject H_0 , blood pressure does not impact outcome

▼ sm.add_constant(x_test, prepend=False)

```

1 x_test = sm.add_constant(x_test, prepend=False)
2 # adding constant to match equation
3 # if we do not add constant, model will follow Binary Logistic Equation but without a constant
4 # if we add constant, model will follow Binary Binary Logistic Binomial Equation with constant
5 x_test.head()
6 # checking columns in Predictor-testing DataFrame after adding constant

```

	Glucose	BloodPressure	Age	const
364	147	74	30	1.0
637	94	76	23	1.0
555	124	70	37	1.0
16	118	84	31	1.0
62	44	62	36	1.0

▼ model.predict(x_test)

```

1 pred_y = mod1.predict(x_test)
2 # generate prediction using x_test sample

```

```

1 pred_y[:5]
2 # printing first 6 records from prediction Series

```

```

364    0.501328
637    0.115016
555    0.388427
16     0.307377
62    -0.137457
dtype: float64

```

```

1 y_test[:5]
2 # printing first 6 records from Response-Testing Series

```

```

364    0
637    0
555    0
16     1
62     0
Name: Outcome, dtype: int64

```

```

1 res = [] # creating empty list to store rounded off prediction values
2 for ctr in pred_y:
3     if ctr < 0.5:
4         res.append(0)
5     else:
6         res.append(1)

```

```

1 res[:5]
2 # printing first 6 records from prediction Series with rounding off

```

```
[0, 0, 0, 0, 1]
```

▼ import confusion_matrix

```

1 from sklearn.metrics import confusion_matrix
2 # importing confusion_matrix module

```

▼ confusion_matrix(y_test, y_pred_rounded)

```

1 confusion_matrix(y_test, res)
2 # generating confusion matrix between Response-Testing Series & Prediction

```

```

array([[12,  0],
       [ 5,  3]])

```

Confusion Matrix

- True / False : Actual/rows
- Positive / Negative : Predicted/columns

Confusion Matrix Structure

		Predicted	
		P	N
Test	T	[TP , TN]	
	F	[FP , FN]	

		P	N
Test		[12 , 0]	
	F	[5 , 3]	

- True Positive TP = 12 [predictions are correct for True]
- True Negative TN = 0 [Type 1 Error]
- False Positive FP = 5 [Type 2 Error]
- False Negative FN = 3 [Predictions are correct for False]
- diagonal elements are erroneous(Type1 or Type2)

Precision & Recall

- Precision = $TP / (TP + FP)$
- Recall = $TP / (TP + FN)$

▼ import classification_report

```
1 from sklearn.metrics import classification_report
2 # importing classification_report
```

▼ classification_report(y_test, y_pred_rounded)

```
1 print(classification_report(y_test, res))
2 # printing classification report
3 # shows precision, recall, f1-score & accuracy
```

	precision	recall	f1-score	support
0	0.71	1.00	0.83	12
1	1.00	0.38	0.55	8
accuracy			0.75	20
macro avg	0.85	0.69	0.69	20
weighted avg	0.82	0.75	0.71	20

```
1 confusion_matrix(y_test, res)
2 # generating confusion matrix between Response-Testing Series & Prediction
```

```
array([[12, 0],
       [ 5, 3]])
```

```
1 12/(12+0)
2 # manually calculating precision
3 # Precision = TP / (TP + FP)
```

```
1.0
```

```
1 12 / (12 + 5)
2 # manually calculating recall
3 # Recall = TP / (TP + FN)
```

```
0.7058823529411765
```

▼ F1-Score

- $F1\text{-Score} = 2 * \text{Precision} * \text{Recall} / (\text{Precision} + \text{Recall})$
- $F1\text{-Score} = 2 * P * R / (P + R)$

```
1 2 * 1.0 * 0.38 / (1.0 + 0.38)
2 # manually calculating F1-Score
```

```
0.5507246376811594
```

```
1 (12 + 3)/(12 + 3 + 0 + 5)
2 # Ratio of correct predictions
```

```
0.75
```

```
1 (1.00 + 0.38)/2
2 # average recall value
```

```
1 (1.00 + 0.38) / (12 + 3 + 0 + 5)
2 # weighted average for recall
```

```
1 (0.71 + 1.00)/2
2 # average precision value
```

```
1 (0.71 + 1.00) / (12 + 3 + 0 + 5)
2 # weighted average for precision
```

▼ Nominal Logistic Regression Model

- when response has multiple categories, and the response does not has a logical order, then we use nominal logistic regression

▼ import statsmodels & sklearn

```
1 import numpy as np
2 import pandas as pd
3 # import numpy & pandas
4
5 import statsmodels
6 import statsmodels.api as sm
7 from statsmodels.formula.api import ols
8 import statsmodels.stats.multicomp
9 # import statsmodels
10
11 import sklearn
12 from sklearn.model_selection import train_test_split
13 from sklearn.metrics import confusion_matrix
14 from sklearn.metrics import classification_report
15 # import sklearn
```

▼ upload dataset

```
1 from google.colab import files
2 uploaded=files.upload()
3 # CDAC_DataBook.xlsx
4 # to be used with google colab
5
```

```

6 # import os
7 # os.chdir(r'C:\Users\surya\Downloads\PG-DBDA-Mar23\Datasets')
8 # os.getcwd()
9 # to change current working directory to specified path
10 # to be used while running on local system

```

▼ `pd.read_excel('WorkBook_Name.xlsx', sheet_name='SheetName')`

```

1 df = pd.read_excel('CDAC_DataBook.xlsx', sheet_name = 'nominal')
2 # reading excel file with specifying the sheet name to load dataset
3 df.head()

```

	ses	write	math	prog
0	1	35	41	1
1	2	33	41	2
2	3	39	44	3
3	1	37	42	1
4	2	31	40	2

```

1 df = df.drop('write', axis=1)
2 # selecting specific/required columns
3 # by dropping unwanted columns from dataset
4 df.head()
5 # checking dataset after selectng only required columns

```

	ses	math	prog
0	1	41	1
1	2	41	2
2	3	44	3
3	1	42	1
4	2	40	2

▼ `pd.get_dummies(categorical_col, drop_first=True)`

```

1 ses_dummy = pd.get_dummies(df['ses'], drop_first=True)
2 # 'ses' is categorical data, so need to create dummies

```

```
3 # creating dummy column
```

```
1 df = df.drop('ses', axis=1)
2 # dropping actual categorical column
```

```
1 df = pd.concat([df, ses_dummy], axis=1)
2 # concatenating dummy column in place of actual categorical column
3 df.head()
4 # checking columns in DataFrame after concatenating categorical column
```

	math	prog	2	3
0	41	1	0	0
1	41	2	1	0
2	44	3	0	1
3	42	1	0	0
4	40	2	1	0

▼ `train_test_split(predictor_cols, response_col, test_size=0.25)`

```
1 x_train, x_test, y_train, y_test = train_test_split(df.drop('prog', axis=1), df['prog'], test_size=0.25)
2 # splitting dataset 4-ways
3 # rows are randomly selected for testing
```

▼ `sm.add_constant(x_train, prepend=False)`

```
1 x_train = sm.add_constant(x_train, prepend=False)
2 # adding constant to match equation
3 # if we do not add constant, model will follow Nominal Logistic Equation but without a constant
4 # if we add constant, model will follow Nominal Logistic Equation with constant
```

▼ `sm.MNLogit(y_train, x_train).fit()`

- uses Multi-Nominal Logistic function

```
1 mod1 = sm.MNLogit(y_train, x_train).fit()
2 # creating model using sm.MNLogit().fit()
```

Optimization terminated successfully.
 Current function value: 0.706854
 Iterations 7

```
1 print(mod1.summary())
2 # printing model summary / MNLogit Regression Results
```

```

=====
MNLogit Regression Results
=====
Dep. Variable:          prog    No. Observations:          150
Model:                  MNLogit  Df Residuals:           142
Method:                  MLE     Df Model:              6
Date:                   Wed, 21 Jun 2023  Pseudo R-squ.:         0.3260
Time:                   06:10:07    Log-Likelihood:        -106.03
converged:               True      LL-Null:              -157.32
Covariance Type:        nonrobust  LLR p-value:          7.239e-20
=====

```

	prog=2	coef	std err	z	P> z	[0.025	0.975]
math	0.0139	0.031	0.445	0.656	-0.047	0.075	
2	2.7011	0.594	4.545	0.000	1.536	3.866	
3	2.4412	1.138	2.145	0.032	0.211	4.671	
const	-1.3353	1.583	-0.844	0.399	-4.437	1.767	

	prog=3	coef	std err	z	P> z	[0.025	0.975]
math	0.1178	0.037	3.189	0.001	0.045	0.190	
2	4.3356	0.947	4.577	0.000	2.479	6.192	
3	6.1418	1.318	4.660	0.000	3.559	8.725	
const	-8.8537	2.250	-3.935	0.000	-13.263	-4.444	

```

=====

```

interpreting rows in MNLogit Regression Summary

- Predictor(X) is ses

```
ses { 1: "Low", 2: "Middle", 3: "High" }
```

- Response(Y) is prog

```
prog { 1: "Vocational", 2: "General", 3: "Academic" }
```

- row1, prog-2: when ses changes from refernce ses 1 to ses 2 then the prob of choosing prog-2
- row2, prog-3: when ses changes from reference ses 1 to ses 3 then the prob of choosing prog-3

▼ interpreting coefficients

- if coefficient is -ve, probability is indicated by movement towards reference value or away from target value
- if coefficient is +ve, probability is indicated by movement towards target value or away from reference value

```
1 np.log(12)
2 # log(n) of n>1 is positive

2.4849066497880004
```

```
1 np.log(0.12)
2 # log(n) of n<1 is negative

-2.120263536200091
```

- in prog-2 row, coeff of 2 = 2.7011 = $\log(\text{prob}(\text{prog2}) / \text{prob}(\text{prog1}))$
- since this value is positive, means $\text{prob}(\text{prog2}) / \text{prob}(\text{prog1}) > 1$ because $\log(n>1)$ is positive
- This further implies $\text{prob}(\text{prog2}) > \text{prob}(\text{prog1})$, because $\text{Numerator/Denominator} > 1$ means $\text{Numerator} > \text{Denominator}$

Establishing Null Hypothesis H_0

- for math column
 - H_0 : math will not impact changing/choosing the prog1 to prog2
- for ses column
 - Surya's H_0 statement: changing ses from reference ses-1 to target ses-2 will not impact the changing the reference prog prog-1 to target prog-2
 - Sudeep's H_0 statement: movement of ses from reference ses-1 to target ses-2 will not influence changing the reference prog-1 to target prog-2

Prediction: math score impacts prog

Q. will score in maths impact the choice of course. if yes, then how?

▼ sm.add_constant(x_test, prepend=False)

```
1 x_test = sm.add_constant(x_test, prepend=False)
2 # adding constant to match equation
```

```
3 # if we do not add constant, model will follow Nominal Logistic Equation but without a constant
4 # if we add constant, model will follow Nominal Logistic Equation with constant
```

▼ model.predict(x_test)

```
1 pred_y = mod1.predict(x_test)
2 # generate prediction using x_test sample
```

```
1 pred_y[:5]
2 # printing first five records from prediction for ses-1, 2, 3
```

	0	1	2
80	0.609625	0.340005	0.050370
119	0.081625	0.632541	0.285834
175	0.580719	0.342423	0.076858
49	0.085070	0.650130	0.264799
85	0.615967	0.338794	0.045239

```
1 x_train.head()
2 # printing first 5 records from Predictor-Training DataFrame
```

	math	2	3	const
148	55	1	0	1.0
88	42	1	0	1.0
139	58	0	1	1.0
135	56	0	1	1.0
89	46	1	0	1.0

```

1 mydata = pd.DataFrame([[60, 1, 0, 1], [70, 1, 0, 1], [80, 1, 0, 1]], columns=['math', '2', '3', 'const'])
2 # manually creating a Predictor-Test DataFrame using random data
3
4 # DataFrame created keeping 'ses' constant as 'ses-2' but varying 'math' col to
5 # see influence of "math" on 'prog' col
6
7 mydata
8 # printing Predictor-Test DataFrame

```

	math	2	3	const
0	60	1	0	1
1	70	1	0	1
2	80	1	0	1

▼ model.predict(Test_DataFrame)

```

1 mod1.predict(mydata)
2 # generating Prediction Series

```

	0	1	2
0	0.043813	0.395686	0.560501
1	0.018896	0.196128	0.784976
2	0.006764	0.080695	0.912541

```

1 mydata = pd.DataFrame([[70, 0, 0, 1], [70, 1, 0, 1], [70, 0, 1, 1]], columns=['math', '2', '3', 'const'])
2 # keeping maths score 70 as constant, and varying ses value
3 # DataFrame created keeping 'math' constant as '70' but varying 'ses' col to
4 # see influence of 'ses' on 'prog' col
5 mydata
6 # printing Predictor-Test DataFrame

```

	math	2	3	const
0	70	0	0	1
1	70	1	0	1
2	70	0	1	1

```
1 mod1.predict(mydata)
2 # generating Prediction Series
```

	0	1	2
0	0.446277	0.310973	0.242749
1	0.018896	0.196128	0.784976
2	0.003819	0.030566	0.965615

▼ Variance Inflation Factor (VIF)

- used to find correlation between two predictors
- VIF should be minimum, $VIF < 10$ is good
- Sometimes, we may accept $VIF \leq 20$, but not more than 20
- If some predictor has $VIF > 20$, then we drop that column to avoid any impact of the high correlation

▼ import statsmodels

```
1 import numpy as np
2 import pandas as pd
3 # import numpy & pandas
4
5 import statsmodels
6 import statsmodels.api as sm
7 from statsmodels.formula.api import ols
8 import statsmodels.stats.multicomp
9 # import statsmodels
10
11 import sklearn
12 from sklearn.model_selection import train_test_split
13 from sklearn.metrics import confusion_matrix
14 from sklearn.metrics import classification_report
15 # import sklearn
```

▼ upload dataset

```
1 from google.colab import files
2 uploaded=files.upload()
3 # CDAC_DataBook.xlsx
4 # to be used with google colab
5
```

```

6 # import os
7 # os.chdir(r'C:\Users\surya\Downloads\PG-DBDA-Mar23\Datasets')
8 # os.getcwd()
9 # to change current working directory to specified path
10 # to be used while running on local system

```

Choose Files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
 Saving CDAC DataBook.xlsx to CDAC DataBook.xlsx

▼ `pd.read_excel('WorkBook_Name.xlsx', sheet_name='SheetName')`

```

1 df = pd.read_excel('CDAC_DataBook.xlsx', sheet_name='VIF')
2 # reading excel file with specifying the sheet name to load dataset
3 df.head()

```

	Education	Region	Gender	Exp	Union	Wage	Age	Race	Occupation	Sector	Married
0	9	0	1	42	0	4.95	57	3	6	1	1
1	12	0	0	1	0	6.67	19	3	6	1	0
2	12	0	0	4	0	4.00	22	3	6	0	0
3	12	0	0	17	0	7.50	35	3	6	0	1
4	13	0	0	9	1	13.07	28	3	6	0	0

▼ `df.columns`

```

1 df.columns
2 # fetching column names

```

```

Index(['Education', 'Region', 'Gender', 'Exp', 'Union', 'Wage', 'Age', 'Race',
      'Occupation', 'Sector', 'Married'],
      dtype='object')

```

```

1 df = df[['Education', 'Gender', 'Exp', 'Age', 'Wage']]
2 # selecting specific columns from dataset

```

▼ `pd.get_dummies(categorical_col, drop_first=True)`

```

1 gd_dummy = pd.get_dummies(df['Gender'], drop_first=True)
2 # 'Gender' is categorical data, so need to create dummies
3 # creating dummy column
4 gd_dummy.head()

```

```

      1
0  1
1  0
2  0
3  0
4  0

```

```

1 df = df.drop('Gender', axis=1)
2 # dropping actual categorical column
3 df.head()

```

	Education	Exp	Age	Wage
0	9	42	57	4.95
1	12	1	19	6.67
2	12	4	22	4.00
3	12	17	35	7.50
4	13	9	28	13.07

```

1 df = pd.concat([df, gd_dummy], axis=1)
2
3 df.head()

```

	Education	Exp	Age	Wage	1
0	9	42	57	4.95	1
1	12	1	19	6.67	0
2	12	4	22	4.00	0
3	12	17	35	7.50	0
4	13	9	28	13.07	0

```
1 x_train = df.drop('Wage', axis=1)
2 y_train = df['Wage']
```

```
1 mod1 = sm.OLS(y_train, x_train).fit()
```

```
1 print(mod1.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          Wage    R-squared (uncentered):          0.817
Model:                  OLS    Adj. R-squared (uncentered):        0.816
Method:                 Least Squares    F-statistic:          591.8
Date:                   Wed, 21 Jun 2023    Prob (F-statistic):      1.10e-193
Time:                   11:26:07    Log-Likelihood:         -1550.9
No. Observations:       533    AIC:                   3110.
Df Residuals:           529    BIC:                   3127.
Df Model:                4
Covariance Type:        nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
Education         1.6271      0.270      6.027      0.000      1.097      2.157
Exp                0.8026      0.205      3.920      0.000      0.400      1.205
Age              -0.6891      0.195     -3.531      0.000     -1.072     -0.306
1                -2.3592      0.388     -6.073      0.000     -3.122     -1.596
=====
Omnibus:                 250.476    Durbin-Watson:           1.870
Prob(Omnibus):           0.000    Jarque-Bera (JB):        2546.360
Skew:                    1.795    Prob(JB):                 0.00
Kurtosis:                13.088    Cond. No.                 94.9
=====

```

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
1 from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
1 variance_inflation_factor(df.values, 0)
```

```
368.9046844691842
```

```
1 variance_inflation_factor(df.values, 1)
```

```
544.6987999131815
```

```
1 variance_inflation_factor(df.values, 2)
```

```
1562.7050138396557
```

```
1 variance_inflation_factor(x_train, 2)
```

```
1526.713272763991
```

- If $0 \leq \text{VIF} \leq 10$, it is okay
- if $10 < \text{VID} \leq 25$, it can be ignored
- for $\text{VIF} > 25$, we need to take action

```
1 x_train.head()
```

	Education	Exp	Age	1
0	9	42	57	1
1	12	1	19	0
2	12	4	22	0
3	12	17	35	0
4	13	9	28	0

```
1 x_train = x_train.drop('Age', axis=1)
```

```
1 mod1 = sm.OLS(y_train, x_train).fit()
```

```
1 print(mod1.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          Wage      R-squared (uncentered):      0.813
Model:                  OLS      Adj. R-squared (uncentered):    0.812
Method:                 Least Squares      F-statistic:          768.2
Date:                   Wed, 21 Jun 2023    Prob (F-statistic):    1.73e-192
Time:                   11:26:27           Log-Likelihood:       -1557.1
No. Observations:       533              AIC:                3120.
Df Residuals:           530              BIC:                3133.
Df Model:                3
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
=====
```



```
-----
Education    0.6777    0.025    27.032    0.000    0.628    0.727
Exp          0.0813    0.014     5.735    0.000    0.053    0.109
1           -2.4805    0.391    -6.342    0.000    -3.249   -1.712
=====
Omnibus:                256.586    Durbin-Watson:                1.815
Prob(Omnibus):           0.000    Jarque-Bera (JB):            2335.399
Skew:                    1.894    Prob(JB):                     0.00
Kurtosis:                12.529    Cond. No.                     48.7
=====
```

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
1 variance_inflation_factor(df.values, 0)
```

```
368.9046844691842
```

```
1 variance_inflation_factor(x_train, 0)
```

```
2.9134749789780923
```

```
1 variance_inflation_factor(x_train, 1)
```

```
2.4837216129295894
```

```
1
```

- if the response is binary, we don't need to create dummy columns

▼ Ordinal Logistic Regression Model

- when response has multiple categories, and the response has a logical order, then we use ordinal logistic regression

▼ import statsmodels

```
1 import numpy as np
2 import pandas as pd
3 # import numpy & pandas
4
5 import statsmodels
6 import statsmodels.api as sm
```

```

7 from statsmodels.formula.api import ols
8 import statsmodels.stats.multicomp
9 # import statsmodels
10
11 import sklearn
12 from sklearn.model_selection import train_test_split
13 from sklearn.metrics import confusion_matrix
14 from sklearn.metrics import classification_report
15 # import sklearn

```

▼ upload dataset

```

1 from google.colab import files
2 uploaded=files.upload()
3 # CDAC_DataBook.xlsx
4 # to be used with google colab
5
6 # import os
7 # os.chdir(r'C:\Users\surya\Downloads\PG-DBDA-Mar23\Datasets')
8 # os.getcwd()
9 # to change current working directory to specified path
10 # to be used while running on local system

```

No file chosen

Upload widget is only available when the cell has been executed in

the current browser session. Please rerun this cell to enable.

Saving CDAC_DataBook.xlsx to CDAC_DataBook.xlsx

▼ import sys

```

1 # import sys
2 # sys.path.append(r'D:/advanced-analytics-files/day10')
3 # to add a path to interpreter for current session
4 # to be used while running on local system
5 # optional : to avoid setting file path each time

```

▼ install mord

```
1 pip install mord
```

Collecting mord

Downloading mord-0.7.tar.gz (8.6 kB)

Preparing metadata (setup.py) ... done

Building wheels for collected packages: mord

```
Building wheel for mord (setup.py) ... done
Created wheel for mord: filename=mord-0.7-py3-none-any.whl size=9885 sha256=e12fe207f39d1f30c7ed935b0ef39a18a044b595a24023fd45c94db8ad199241
Stored in directory: /root/.cache/pip/wheels/77/00/19/3cea86fbfc737ec4acb515cd94497dcc33f943fa157548b96c
Successfully built mord
Installing collected packages: mord
Successfully installed mord-0.7
```

▼ import mord.LogisticAT

```
1 import mord
2 from mord import LogisticAT
```

▼ pd.read_excel('WorkBook_Name.xlsx', sheet_name='SheetName')

```
1 df = pd.read_excel('CDAC_DataBook.xlsx', sheet_name='ordinal')
2 # reading excel file with specifying the sheet name to load dataset
3 df.head()
```

	Survival	Region	ToxicLevel
0	1	1	62.0
1	1	2	46.0
2	2	1	48.5
3	3	2	32.0
4	2	1	63.5

▼ pd.get_dummies(categorical_col, drop_first=True)

```
1 reg_dummy = pd.get_dummies(df['Region'], drop_first=True)
2 # 'Region' is categorical data, either 1 or 2, so need to create dummies
3 # creating dummy column
4 reg_dummy.head()
5
```

2

0 0

1 1

```
1 df = df.drop('Region', axis=1)
2 # dropping actual categorical column 'Region'
3 df.head()
```

	Survival	ToxicLevel
0	1	62.0
1	1	46.0
2	2	48.5
3	3	32.0
4	2	63.5

```
1 df = pd.concat([df, reg_dummy], axis=1)
2 # concatenating dummy column in place of actual categorical column
3 df.head()
4 # checking columns in DataFrame after concatenating categorical column
```

	Survival	ToxicLevel	2
0	1	62.0	0
1	1	46.0	1
2	2	48.5	0
3	3	32.0	1
4	2	63.5	0

▼ train_test_split(predictor_cols, response_col, test_size=0.25)

```
1 x_train, x_test, y_train, y_test = train_test_split(df.drop('Survival', axis=1), df['Survival'], test_size=0.25)
2 # splitting dataset 4-ways
3 # rows are randomly selected for testing
```

▼ sm.add_constant(x_train, prepend=False)

```

1 x_train = sm.add_constant(x_train, prepend=False)
2 # adding constant to match equation
3 # if we do not add constant, model will follow Nominal Ordinal Equation but without a constant
4 # if we add constant, model will follow Nominal Ordinal Equation with constant

```

▼ mord.LogisticAT().fit(x_train, y_train)

```

1 mod1 = LogisticAT().fit(x_train, y_train)
2 # creating model using mord.LogisticAT().fit(x_train, y_train)

```

```

1 # print(mod1.summary())
2 # cannot print summary for LogisticAT() model
3 # as it does not have summary() method

```

```
<bound method LogisticAT.score of LogisticAT()>
```

```

1 x_test = sm.add_constant(x_test, prepend=False)
2 # adding constant to match equation
3 # if we do not add constant, model will follow Nominal Ordinal Equation but without a constant
4 # if we add constant, model will follow Nominal Ordinal Equation with constant

```

▼ model.predict(x_test)

```

1 pred_y = mod1.predict(x_test)
2 # generate prediction using x_test sample
3 pred_y[:5]

```

```
array([2, 2, 2, 1, 2])
```

▼ confusion_matrix(y_test, y_pred_rounded)

```

1 print(confusion_matrix(y_test, pred_y))
2 # generating confusion matrix between Response-Testing Series & Prediction

```

```

[[ 0  4  0]
 [ 2 10  0]
 [ 0  3  0]]

```

```
1
```

Double-click (or enter) to edit

▼ Counts Regression Model

1. Poisson Regression Model
2. Negative Binomial Regression Model

- uses `.from_formula(Resp ~ P1 + P2 + P2)` to represent the relation between response & predictors

▼ Poisson Regression

- when response is discrete data
- when the variation is expected to be low

▼ import statsmodels

```
1 import numpy as np
2 import pandas as pd
3 # import numpy & pandas
4
5 import statsmodels
6 import statsmodels.api as sm
7 from statsmodels.formula.api import ols
8 import statsmodels.stats.multicomp
9 # import statsmodels
10
11 import sklearn
12 from sklearn.model_selection import train_test_split
13 from sklearn.metrics import confusion_matrix
14 from sklearn.metrics import classification_report
15 # import sklearn
```

▼ upload dataset

```
1 from google.colab import files
2 uploaded=files.upload()
3 # CDAC_DataBook.xlsx
4 # to be used with google colab
5
6 # import os
7 # os.chdir(r'C:\Users\surya\Downloads\PG-DBDA-Mar23\Datasets')
```

```

8 # os.getcwd()
9 # to change current working directory to specified path
10 # to be used while running on local system

```

▼ `pd.read_excel('WorkBook_Name.xlsx', sheet_name='SheetName')`

```

1 df = pd.read_excel('CDAC_DataBook.xlsx', sheet_name='poisson')
2 # reading excel file with specifying the sheet name to load dataset
3 df.head()

```

	num_awards	prog	math
0	0	3	41
1	0	1	41
2	0	3	44
3	0	3	42
4	0	3	40

▼ `pd.get_dummies(categorical_col, drop_first=True)`

```

1 prog_dummy = pd.get_dummies(df['prog'], drop_first=True)
2 # prog is categorical data, so need to create dummies
3 # creating dummy column
4 prog_dummy.head()

```

	2	3
0	0	1
1	0	0
2	0	1
3	0	1
4	0	1

```

1 df = df.drop('prog', axis=1)
2 # dropping actual categorical column
3 df.head()
4 # checking columns in DataFrame after dropping categorical column

```

	num_awards	math
0	0	41
1	0	41
2	0	44
3	0	42
4	0	40

```
1 df = pd.concat([df, prog_dummy], axis=1)
2 # concatenating dummy column in place of actual categorical column
3 df.head()
4 # checking columns in DataFrame after concatenating categorical column
```

	num_awards	math	2	3
0	0	41	0	1
1	0	41	0	0
2	0	44	0	1
3	0	42	0	1
4	0	40	0	1

▼ train_test_split(predictor_cols, response_col, test_size=0.25)

```
1 x_train, x_test, y_train, y_test = train_test_split(df.drop('num_awards', axis=1), df['num_awards'], test_size=0.25)
2 # splitting dataset 4-ways
3 # rows are randomly selected for testing
```

▼ sm.add_constant(x_train, prepend=False)

```
1 x_train = sm.add_constant(x_train, prepend=False)
2 # adding constant to match equation
3 # if we do not add constant, model will follow Poisson Equation but without a constant
4 # if we add constant, model will follow Poisson Equation with constant
5 x_train.head()
6 # checking columns in Predictor-training DataFrame after adding constant
```

▼ import Poisson


```
1 from statsmodels.discrete.discrete_model import Poisson as psn
2 # import statsmodels.discrete.discrete_model.Poisson
```

▼ `pd.concat([x_train, y_train], axis=1)`

```
1 df_train = pd.concat([x_train, y_train], axis=1)
2 # # concatenating Predictor(X) & Response(Y) into single DataFrame
3 # so that relation can be established in formula
4 df_train.head()
5 # checking columns in concatenated training DataFrame
```

	math	2	3	const	num_awards
148	55	1	0	1.0	0
199	73	1	0	1.0	3
27	46	0	0	1.0	1
93	50	1	0	1.0	0
118	54	1	0	1.0	0

▼ `df_train.columns`

```
1 x = df_train.columns
2 # fetching column names
3 x
4 # printing column names
```

```
Index(['math', 2, 3, 'const', 'num_awards'], dtype='object')
```

▼ `df_train.rename(columns={2:'Col2_name', 3:'Col3_name'}, inplace=True)`

```
1 df_train.rename(columns={2:'prog2', 3:'prog3'}, inplace=True)
2 # rename columns because column name is appearing as column index
3 df_train.head()
4 # checking columns in renamed training DataFrame
```

	math	prog2	prog3	const	num_awards
148	55	1	0	1.0	0
199	73	1	0	1.0	3
27	46	0	0	1.0	1
...

▼ `Poisson.from_formula('Response ~ P1 + p2 + p3', data=DataSet).fit()`

```
1 mod1 = psn.from_formula('num_awards ~ math + prog2 + prog3', data=df_train).fit()
2 # establishing relation in formula
3 # creating model using Poisson.from_formula() method
```

```
Optimization terminated successfully.
Current function value: 0.990968
Iterations 6
```

```
1 print(mod1.summary())
2 # printing model summary / Poisson Regression Results
```

```

Poisson Regression Results
=====
Dep. Variable:          num_awards    No. Observations:          150
Model:                  Poisson      Df Residuals:              146
Method:                  MLE         Df Model:                  3
Date:                   Wed, 21 Jun 2023    Pseudo R-squ.:            0.2003
Time:                   12:25:29          Log-Likelihood:           -148.65
converged:              True            LL-Null:                  -185.87
Covariance Type:        nonrobust        LLR p-value:              4.764e-16
=====
              coef    std err          z      P>|z|      [0.025    0.975]
-----
Intercept    -4.9471      0.742     -6.667    0.000     -6.402     -3.493
math          0.0699      0.012      5.634    0.000      0.046      0.094
prog2         0.8920      0.364      2.452    0.014      0.179      1.605
prog3         0.1469      0.487      0.301    0.763     -0.808      1.102
=====
```

- H_0 : if prog changes prog-1 \rightarrow prog-2 , then the number of awards is not impacted
 - since P-Value $0.014 < 0.05$, so we reject this H_0 , means if prog changes prog-1 \rightarrow prog-2 , then number of awards changes
- H_0 : if prog changes prog-1 \rightarrow prog-3 , then the number of awards is not impacted
 - since P-Value $0.763 > 0.05$, so we do not reject this H_0 , means if prog changes prog-1 \rightarrow prog-3 , then number of awards does not change

1

▼ Negative Binomial Regression Model

- when response is discrete data
- when the variation is expected to be high

▼ import statsmodels

```

1 import numpy as np
2 import pandas as pd
3 # import numpy & pandas
4
5 import statsmodels
6 import statsmodels.api as sm
7 from statsmodels.formula.api import ols
8 import statsmodels.stats.multicomp
9 # import statsmodels
10
11 import sklearn
12 from sklearn.model_selection import train_test_split
13 from sklearn.metrics import confusion_matrix
14 from sklearn.metrics import classification_report
15 # import sklearn

```

▼ upload dataset

```

1 from google.colab import files
2 uploaded=files.upload()
3 # CDAC_DataBook.xlsx
4 # to be used with google colab
5
6 # import os

```

```

7 # os.chdir(r'C:\Users\surya\Downloads\PG-DBDA-Mar23\Datasets')
8 # os.getcwd()
9 # to change current working directory to specified path
10 # to be used while running on local system

```

[Choose Files](#) No file chosen

Upload widget is only available when the cell has been executed in

the current browser session. Please rerun this cell to enable.

Saving CDAC DataBook.xlsx to CDAC DataBook.xlsx

▼ `pd.read_excel('WorkBook_Name.xlsx', sheet_name='SheetName')`

```

1 df = pd.read_excel('CDAC_DataBook.xlsx', sheet_name='neg_bin')
2 # reading excelfile with specifying the sheet name to load dataset
3 df.head()

```

	math	prog	daysabs
0	63	Academic	4
1	27	Academic	4
2	20	Academic	2
3	16	Academic	3
4	2	Academic	3

▼ `import NegativeBinomial`

```

1 from statsmodels.discrete.discrete_model import NegativeBinomial as ngb
2 # import statsmodels.discrete.discrete_model.NegativeBinomial

```

▼ `pd.get_dummies(categorical_col, drop_first=True)`

- create dummies for categorical columns
- first value in alphabetical order is automatically taken as reference value

```

1 prog_dummy = pd.get_dummies(df['prog'], drop_first=True)
2 # creating dummy column
3 prog_dummy.head()

```

	General	Vocational
0	0	0
1	0	0
2	0	0
3	0	0

```

1 df = df.drop('prog', axis=1)
2 # dropping actual categorical column
3 df.head()
4 # checking columns in DataFrame after dropping categorical column

```

	math	daysabs
0	63	4
1	27	4
2	20	2
3	16	3
4	2	3

```

1 df = pd.concat([df, prog_dummy], axis=1)
2 # concatenating dummy column in place of actual categorical column
3 df.head()
4 # checking columns in DataFrame after concatenating categorical column

```

	math	daysabs	General	Vocational
0	63	4	0	0
1	27	4	0	0
2	20	2	0	0
3	16	3	0	0
4	2	3	0	0

▼ train_test_split(predictor_cols, response_col, test_size=0.25)

```
1 x_train, x_test, y_train, y_test = train_test_split(df.drop('daysabs', axis=1), df['daysabs'], test_size=0.25)
2 # splitting dataset 4-ways
3 # rows are randomly selected for testing
```

▼ `sm.add_constant(x_train, prepend=False)`

```
1 x_train = sm.add_constant(x_train, prepend=False)
2 # adding constant to match equation
3 # if we do not add constant, model will follow Negative Binomial Equation but without a constant
4 # if we add constant, model will follow Binary Negative Binomial Equation with constant
5 x_train.head()
6 # checking columns in Predictor-training DataFrame after adding constant
```

	math	General	Vocational	const
182	24	0	1	1.0
202	31	0	1	1.0
188	65	0	0	1.0
197	57	0	0	1.0
148	23	0	0	1.0

```
1 y_train.head()
2 # checking columns in Response-training Series
```

```
182    1
202    1
188    0
197    0
148    12
Name: daysabs, dtype: int64
```

▼ `pd.concat([x_train, y_train], axis=1)`

```
1 df_train = pd.concat([x_train, y_train], axis=1)
2 # concatenating Predictor(X) & Response(Y) into single DataFrame
3 # so that relation can be established in formula
4 df_train.head()
5 # checking columns in concatenated training DataFrame
```

	math	General	Vocational	const	daysabs
182	24	0	1	1.0	1
202	31	0	1	1.0	1
188	65	0	0	1.0	0
197	57	0	0	1.0	0

▼ `NegativeBinomial.from_formula('Response ~ P1 + P2 + P3', data=DataSet).fit()`

```
1 mod1 = nbg.from_formula('daysabs ~ math + General + Vocational', data=df_train).fit()
2 # establishing relation in formula
3 # creating model using NegativeBinomial.from_formula() method
```

```
Optimization terminated successfully.
Current function value: 2.745232
Iterations: 16
Function evaluations: 23
Gradient evaluations: 23
```

```
1 print(mod1.summary())
2 # printing model summary / NegativeBinomial Regression Results
```

```

                NegativeBinomial Regression Results
=====
Dep. Variable:          daysabs    No. Observations:          235
Model:                NegativeBinomial    Df Residuals:          231
Method:                  MLE    Df Model:              3
Date:                  Thu, 29 Jun 2023    Pseudo R-squ.:        0.04603
Time:                  10:10:07    Log-Likelihood:       -645.13
converged:              True    LL-Null:             -676.25
Covariance Type:        nonrobust    LLR p-value:         1.942e-13
=====
              coef    std err          z      P>|z|      [0.025    0.975]
-----
Intercept      2.2312      0.151     14.812     0.000      1.936      2.526
math           -0.0057      0.003     -2.058     0.040     -0.011     -0.000
General         0.4778      0.209      2.282     0.023      0.067      0.888
Vocational     -1.0546      0.164     -6.418     0.000     -1.377     -0.733
alpha           0.9057      0.109      8.328     0.000      0.693      1.119
=====
```

Establishing Null Hypothesis H_0

- H_0 : maths(predictor) does not influence the daysabs(response)
- H_A : maths(predictor) influences the daysabs(response)

- for maths- P-Value $0.040 < 0.05$, so we reject the H_0 , means maths score will impact the daysabs

interpreting coefficients

- sign of coefficient being -ve means, response will decrease if predictor increases.
 - for maths score, coefficient is -ve means, if maths score increases, daysabs decrease
- impact of days of absence in General > days of absence in Academic
- impact of days of absence in vocational < days of absence in Academic
- means highest absent days in General, lowest absent days in Vocational

▼ sm.add_constant(x_test, prepend=False)

```
1 x_test = sm.add_constant(x_test, prepend=False)
2 # adding constant to match equation
3 # if we do not add constant, model will follow Negative Binomial Equation but without a constant
4 # if we add constant, model will follow Binary Negative Binomial Equation with constant
5 x_test.head()
6 # checking columns in Predictor-testing DataFrame after adding constant
```

