

▼ Classification using Logistic Regression & K-NN Classifier

▼ import libs

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

▼ import dataset

```
1 # from google.colab import files
2 # uploaded = files.upload()
3 # D8data1.csv
4
5 import os
6 os.chdir(r'C:\Users\surya\Downloads\PG-DBDA-Mar23\Datasets')
7 os.getcwd()
```

📁 'C:\\Users\\surya\\Downloads\\PG-DBDA-Mar23\\Datasets'

```
1 data = pd.read_csv('D8data1.csv')
2 data.head()
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
1 data.shape
```

(400, 5)

```
1 data.describe()
```

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

```
1 data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User ID         400 non-null   int64
1   Gender          400 non-null   object
2   Age             400 non-null   int64
3   EstimatedSalary 400 non-null   int64
4   Purchased       400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

▼ EDA

▼ count for Gender column

```
1 sns.countplot(x='Purchased', hue='Gender', data=data)
```

<Axes: xlabel='Purchased', ylabel='count'>




```
1 # alt for countplot
2 # sns.countplot(x=data['Purchased'], hue=data['Gender'])
```

Purchased

▼ dist for Age column

```
1 sns.displot(data['Age'])
```

```
c:\users\surya\appdata\local\programs\python\python39\lib\site-packages\seaborn\axisgrid.p  
self._figure.tight_layout(*args, **kwargs)  
<seaborn.axisgrid.FacetGrid at 0x222bd492340>
```




```
1 sns.distplot(data['Age'])
```

▼ dist for EstimatedSalary column

```
1 data['EstimatedSalary'].plot(kind='kde')  
2 # kernel density plot
```

<Axes: ylabel='Density'>



▼ identify X & Y

```
1 x = data[['Age', 'EstimatedSalary']]  
2 x[:5]
```

	Age	EstimatedSalary
0	19	19000
1	0	
2	0	
3	0	
4	0	

```

1 y = data['Purchased']
2 y[:5]

0    0
1    0
2    0
3    0
4    0
Name: Purchased, dtype: int64

```

▼ Splitting

```
1 from sklearn.model_selection import train_test_split
```

```
1 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=0)
```

```
1 x_train.head()
```

	Age	EstimatedSalary
250	44	39000
63	32	120000
312	38	50000
159	32	135000
283	52	21000

```
1 x_test.head()
```

Age EstimatedSalary

```
1 x_train.describe()
2 # to get an idea of min & max for scaling values
```

	Age	EstimatedSalary
count	300.000000	300.000000
mean	38.126667	69583.333333
std	10.114592	34548.541619
min	18.000000	15000.000000
25%	30.750000	43000.000000
50%	38.000000	69500.000000
75%	46.000000	88000.000000
max	60.000000	150000.000000

▼ Preprocessing

▼ Feature Scaling

```
1 from sklearn.preprocessing import StandardScaler
```

```
1 scaler = StandardScaler()
```

```
1 x_train = scaler.fit_transform(x_train)
2 x_train[:5]
```

```
array([[ 0.58164944, -0.88670699],
       [-0.60673761,  1.46173768],
       [-0.01254409, -0.5677824 ],
       [-0.60673761,  1.89663484],
       [ 1.37390747, -1.40858358]])
```

```
1 x_test = scaler.fit_transform(x_test)
2 x_test[:5]
```

```
array([[ -0.54748976,  0.5130727 ],
       [ 0.15442019, -0.61825566],
       [-0.10879604,  0.14615539],
       [-0.54748976,  0.26846116],
       [-0.10879604, -0.61825566]])
```

▼ Logistic Regression Classification

```
1 from sklearn.linear_model import LogisticRegression
```

▼ Modeling

```
1 model = LogisticRegression()
```

▼ Training

```
1 model.fit(x_train, y_train)
```

```
▼ LogisticRegression
LogisticRegression()
```

▼ Prediction

```
1 prediction = model.predict(x_test)
2 # generating prediction
3 prediction[:5]
```

```
array([0, 0, 0, 0, 0], dtype=int64)
```

▼ Evaluation

- checking performance using evaluation metrics

▼ confusion_matrix


```
1 from sklearn.metrics import confusion_matrix
```

```
1 confusion_matrix(y_test, prediction)
```

```
array([[63,  5],
       [ 8, 24]], dtype=int64)
```

▼ classification_report

```
1 from sklearn.metrics import classification_report
```

```
1 print(classification_report(y_test, prediction))
```

	precision	recall	f1-score	support
0	0.89	0.93	0.91	68
1	0.83	0.75	0.79	32
accuracy			0.87	100
macro avg	0.86	0.84	0.85	100
weighted avg	0.87	0.87	0.87	100

▼ accuracy_score

```
1 from sklearn.metrics import accuracy_score
```

```
1 accuracy_score(y_test, prediction)
```

```
0.87
```

▼ precision_score

```
1 from sklearn.metrics import precision_score
```

```
1 precision_score(y_test, prediction)
```

```
0.8275862068965517
```

▼ recall_score

```
1 from sklearn.metrics import recall_score
```

```
1 recall_score(y_test, prediction)
```

```
0.75
```

▼ K-Nearest Neighbor (KNN) Classification

▼ Modeling

```
1 from sklearn.neighbors import KNeighborsClassifier
```

```
1 knn_model = KNeighborsClassifier(n_neighbors=5)
```

▼ Training

```
1 knn_model.fit(x_train, y_train)
```

```
▼ KNeighborsClassifier  
KNeighborsClassifier()
```

▼ Prediction

```
1 knn_predictions = knn_model.predict(x_test)  
2 knn_predictions[:5]
```

```
array([0, 0, 0, 0, 0], dtype=int64)
```

▼ Evaluation

▼ confusion_matrix

```
1 from sklearn.metrics import confusion_matrix
```

```
1 confusion_matrix(y_test, knn_predictions)
```

```
array([[64,  4],
       [ 3, 29]], dtype=int64)
```

▼ classification_report

```
1 from sklearn.metrics import classification_report
```

```
1 print(classification_report(y_test, knn_predictions))
```

	precision	recall	f1-score	support
0	0.96	0.94	0.95	68
1	0.88	0.91	0.89	32
accuracy			0.93	100
macro avg	0.92	0.92	0.92	100
weighted avg	0.93	0.93	0.93	100

▼ accuracy_score

```
1 from sklearn.metrics import accuracy_score
```

```
1 accuracy_score(y_test, knn_predictions)
```

```
0.93
```

▼ precision_score

```
1 from sklearn.metrics import precision_score
```

```
1 precision_score(y_test, knn_predictions)
```

0.8787878787878788

▼ recall_score

```
1 from sklearn.metrics import recall_score
```

```
1 recall_score(y_test, knn_predictions)
```

0.90625

Observation Note

- Note that K-NN model has more accuracy than Logistic, so K-NN performs better than Logistic regression

▼ HW:

- conditional probability
- Bayes Theorem

1

