

## ▼ Support Vector Machines (SVM)

- Support vector machines (SVMs) are a set of supervised learning methods used for classification, regression and outliers detection
- Goal is to create the best line or decision boundary called optimal hyperplane that can segregate n-dimensional space into classes
- SVM finds the hyperplane using `support vectors` (essential training tuples) and `margins` (defined by the support vectors)
- Support Vector Machines tries to produce linear decision boundaries
- is robust to outliers
- based on statistical approach
- can better handle highly-dimensional data
- used for classification as well as Regression problems, but mostly used for Classification
- is a Classifier, forward neural network, supervised learning algorithm
- Can handle linear as well as non-linear data
- SVM Algorithm types
  - Linear SVM
    - for linearly separable data
    - a single straight line can entirely divide the data points into their respective classes
  - Kernel SVM
    - for non-linear separable data
    - used when data points cannot be separated with a single straight line
    - original data is transformed by these kernel functions into a higher dimensional feature space where the features can be linearly separable

- **Hyperplane**
  - the best decision boundary differentiating classes
  - it can be linear/straight line for linearly separable data / 2 features
  - it can be non-linear plane / 2-D plane for non-linearly separable data / 3 features
  - equation is  $w^T x + b = 0$
- **Optimal Hyperplane**
  - the hyperplane with maximum margin is called the optimal Hyperplane
- **Support Vectors**
  - vectors / data points closest to the hyperplane are called Support Vectors
- **Margin**
  - distance between the hyperplane and the support vectors (nearest data point) is called Margin
  - margin is expected to be as large as possible to find the optimal hyperplane
  - two types of margins
    1. Hard Margin
      - **Maximum Margin Hyperplane Or Hard Margin Hyperplane** is a hyperplane that properly separates the data points of different categories without any misclassification
    2. Soft Margin
      - when data is not perfectly separable or contains outliers, then Soft margin is used
      - Each data point has a slack variable introduced by the soft margin SVM formulation, which softens the strict margin requirement and permits certain misclassifications or violations
      - it discovers compromise between increasing the margin and reducing violations
- **Kernel**
  - mathematical function in SVM used to map the original input data points to high dimensional feature spaces
  - makes it easy to find the hyperplane even if the data points are not linearly separable

- Kernel refers to a method that allows us to apply linear classifiers to non-linear problems by mapping non-linear data into a higher-dimensional space without the need to visit or understand that higher-dimensional space
- Common kernel functions are linear, polynomial, rbf(radial basis function), sigmoid, neural net
- `C`
  - called as regularization parameter used to balance margin maximization and misclassification fines
  - penalty for misclassifications is decided by regularization parameter
  - more the value of `C` , stricter the penalty, so leading to smaller margin and fewer misclassifications
- Model performance can be altered by changing the value of hyperparameters which are `C` (Regularization factor), `gamma`, and `kernel`

## Advantages of SVM

- works better when data is linear, effective in high dimensional data
- robust to outliers
- can help us with image classification
- memory efficient as it uses subset of training data points in the decision function called support vectors
- different kernel functions can be specified and it is possible to specify custom kernels

## Disadvantages of SVM

- Choosing a good kernel is not easy
- does not show good results on a bigger data set
- Hyperparameters of SVM are `C` and `gamma` , & it is not easy to fine tune these hyper parameters

## ▼ Note for Kernel Trick

- when there is no separating plane
  - Use bigger set of features, makes use of kernel trick
    - it would make computation hopelessly slow, but using kernel trick we can make computation fast even with huge number of features
  - extend the definition of maximum margin to allow non-separating planes
    - this can be done by using the "Slack" variables, thereby using soft margin technique
    - slack variables are constrained to be non-negative

## ▼ import libs

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

## ▼ import dataset

```
1 # from google.colab import files
2 # uploaded = files.upload()
3 # D9data3.csv
4
5 import os
6 os.chdir(r'C:\Users\surya\Downloads\PG-DBDA-Mar23\Datasets')
7 os.getcwd()
```

```
'C:\\Users\\surya\\Downloads\\PG-DBDA-Mar23\\Datasets '
```

```
1 dataset = pd.read_csv('D9data3.csv')
2 dataset.head()
```

	Age	EstimatedSalary	Purchased
<b>0</b>	19	19000	0
<b>1</b>	35	20000	0
<b>2</b>	26	43000	0
<b>3</b>	27	57000	0
<b>4</b>	19	76000	0

```
1 dataset.shape
```

```
(400, 3)
```

```
1 dataset.describe()
```

	Age	EstimatedSalary	Purchased
<b>count</b>	400.000000	400.000000	400.000000
<b>mean</b>	37.655000	69742.500000	0.357500
<b>std</b>	10.482877	34096.960282	0.479864
<b>min</b>	18.000000	15000.000000	0.000000
<b>25%</b>	29.750000	43000.000000	0.000000
<b>50%</b>	37.000000	70000.000000	0.000000
<b>75%</b>	46.000000	88000.000000	1.000000
<b>max</b>	60.000000	150000.000000	1.000000

```
1 dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              400 non-null    int64
1   EstimatedSalary  400 non-null    int64
2   Purchased        400 non-null    int64
dtypes: int64(3)
memory usage: 9.5 KB
```

## ▼ Imputation (Null check)

```
1 dataset.isnull().sum()
```

```
Age              0
EstimatedSalary  0
Purchased        0
dtype: int64
```

## ▼ identify X & Y

```
1 x = dataset.iloc[ : , :-1].values
2 x[:5]
```

```
array([[ 19, 19000],
       [ 35, 20000],
       [ 26, 43000],
       [ 27, 57000],
       [ 19, 76000]], dtype=int64)
```

```
1 y = dataset.iloc[ : , -1].values
2 y[:5]

array([0, 0, 0, 0, 0], dtype=int64)
```

## ▼ splitting

```
1 from sklearn.model_selection import train_test_split
```

```
1 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)
```

```
1 x_train[:5]

array([[ 58, 144000],
       [ 59,  83000],
       [ 24,  55000],
       [ 26,  35000],
       [ 58,  38000]], dtype=int64)
```

```
1 y_train[:5]

array([1, 0, 0, 0, 1], dtype=int64)
```

## ▼ Preprocessing

### ▼ Feature scaling

```
1 from sklearn.preprocessing import StandardScaler
```

```
1 sc = StandardScaler()
```

```
1 x_train = sc.fit_transform(x_train)
```

```
2 x_train[:5]
```

```
array([[ 1.92295008,  2.14601566],  
       [ 2.02016082,  0.3787193 ],  
       [-1.3822153 , -0.4324987 ],  
       [-1.18779381, -1.01194013],  
       [ 1.92295008, -0.92502392]])
```

```
1 x_test = sc.fit_transform(x_test)
```

```
2 x_test[:5]
```

```
array([[ -0.49618606,  0.56021375],  
       [ 0.2389044 , -0.59133674],  
       [-0.03675452,  0.18673792],  
       [-0.49618606,  0.31122986],  
       [-0.03675452, -0.59133674]])
```

## ▼ Linear SVM

- for linearly separable data

## ▼ Modeling - Linear SVM

```
1 from sklearn.svm import SVC
```

```
1 lsvmclassifier = SVC(C=1, kernel='linear', random_state=0)
```

```
2 # C-Support Vector Classification
```

```
3 # C : Regularization parameter
```



## ▼ Training - Linear SVM

```
1 lsvmclassifier.fit(x_train, y_train)
```

```
▼ SVC
SVC(C=1, kernel='linear', random_state=0)
```

## ▼ Prediction - Linear SVM

```
1 lsvmclassifier.predict(sc.transform([[30, 78000]]))
2 # prediction with custom test case
```

```
array([0], dtype=int64)
```

```
1 y_pred_linear = lsvmclassifier.predict(x_test)
2 y_pred_linear[:5]
```

```
array([0, 0, 0, 0, 0], dtype=int64)
```

## ▼ Evaluation - Linear SVM

### ▼ confusion\_matrix

```
1 from sklearn.metrics import confusion_matrix
```

```
1 confusion_matrix(y_test, y_pred_linear)
```

```
array([[52,  6],
       [ 3, 19]], dtype=int64)
```

### ▼ classification\_report

```
1 from sklearn.metrics import classification_report
```

```
1 print(classification_report(y_test, y_pred_linear))
```

	precision	recall	f1-score	support
0	0.95	0.90	0.92	58
1	0.76	0.86	0.81	22
accuracy			0.89	80
macro avg	0.85	0.88	0.86	80
weighted avg	0.89	0.89	0.89	80

### ▼ accuracy\_score

```
1 from sklearn.metrics import accuracy_score
```

```
1 accuracy_score(y_test, y_pred_linear)
```

```
0.8875
```

### ▼ precision\_score

```
1 from sklearn.metrics import precision_score
```

```
1 precision_score(y_test, y_pred_linear)
```

0.76

### ▼ recall\_score

```
1 from sklearn.metrics import recall_score
```

```
1 recall_score(y_test, y_pred_linear)
```

0.8636363636363636

### ▼ Kernel SVM (c=1)

- for non-linearly separable data

### ▼ Modeling - Kernel SVM (c=1)

```
1 from sklearn.svm import SVC
```

```
1 ksvmclassifier = SVC(C=1, kernel='rbf', random_state=0)
```

### ▼ Training - Kernel SVM (c=1)

```
1 ksvmclassifier.fit(x_train, y_train)
```

▼ SVC

SVC(C=1, random\_state=0)

## ▼ Prediction - Kernel SVM (c=1)

```
1 y_pred_kernel = ksvmclassifier.predict(x_test)
2 y_pred_kernel[:5]
```

```
array([0, 0, 0, 0, 0], dtype=int64)
```

## ▼ Evaluation Kernel SVM (c=1)

### ▼ confusion\_matrix

```
1 from sklearn.metrics import confusion_matrix
```

```
1 confusion_matrix(y_test, y_pred_kernel)
```

```
array([[54, 4],
       [ 1, 21]], dtype=int64)
```

### ▼ classification\_report

```
1 from sklearn.metrics import classification_report
```

```
1 print(classification_report(y_test, y_pred_kernel))
```

	precision	recall	f1-score	support
0	0.98	0.93	0.96	58

1	0.84	0.95	0.89	22
accuracy			0.94	80
macro avg	0.91	0.94	0.92	80
weighted avg	0.94	0.94	0.94	80

### ▼ accuracy\_score

```
1 from sklearn.metrics import accuracy_score
```

```
1 accuracy_score(y_test, y_pred_kernel)
2 # higher accuracy means mis-classification has reduced
```

```
0.9375
```

### ▼ precision\_score

```
1 from sklearn.metrics import precision_score
```

```
1 precision_score(y_test, y_pred_kernel)
```

```
0.84
```

### ▼ recall\_score

```
1 from sklearn.metrics import recall_score
```

```
1 recall_score(y_test, y_pred_kernel)
```

0.9545454545454546

## ▼ Kernel SVM (c=100)

## ▼ Modeling - Kernel SVM (c=100)

```
1 from sklearn.svm import SVC
```

```
1 ksvmclassifier100 = SVC(C=100, kernel='rbf', random_state=0)
```

## ▼ Training - Kernel SVM (c=100)

```
1 ksvmclassifier100.fit(x_train, y_train)
```

▼ SVC  
SVC(C=100, random\_state=0)

## ▼ Prediction - Kernel SVM(c=100)

```
1 y_pred_kernel100 = ksvmclassifier100.predict(x_test)
2 y_pred_kernel100[:5]
```

```
array([0, 0, 0, 0, 0], dtype=int64)
```

## ▼ Evaluation - Kernel SVM (c=100)

### ▼ confusion\_matrix

```
1 from sklearn.metrics import confusion_matrix
```

```
1 confusion_matrix(y_test, y_pred_kernel100)
```

```
array([[54,  4],
       [ 1, 21]], dtype=int64)
```

### ▼ classification\_report

```
1 from sklearn.metrics import classification_report
```

```
1 print(classification_report(y_test, y_pred_kernel100))
```

	precision	recall	f1-score	support
0	0.98	0.93	0.96	58
1	0.84	0.95	0.89	22
accuracy			0.94	80
macro avg	0.91	0.94	0.92	80
weighted avg	0.94	0.94	0.94	80

### ▼ accuracy\_score

```
1 from sklearn.metrics import accuracy_score
```

```
1 accuracy_score(y_test, y_pred_kernel100)
```

```
0.9375
```

### ▼ precision\_score

```
1 from sklearn.metrics import precision_score
```

```
1 precision_score(y_test, y_pred_kernel100)
```

```
0.84
```

### ▼ recall\_score

```
1 from sklearn.metrics import recall_score
```

```
1 recall_score(y_test, y_pred_kernel100)
```

```
0.9545454545454546
```

## ▼ SVM Application

### ▼ import libs



```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
```

## ▼ import dataset

```
1 # from google.colab import files
2 # uploaded = files.upload()
3 # D13data1.csv
4
5 import os
6 os.chdir(r'C:\Users\surya\Downloads\PG-DBDA-Mar23\Datasets')
7 os.getcwd()
```

'C:\\Users\\surya\\Downloads\\PG-DBDA-Mar23\\Datasets'

```
1 dataset = pd.read_csv('D13data1.csv')
2 dataset.head()
```

	Mean of the integrated profile	Standard deviation of the integrated profile	Excess kurtosis of the integrated profile	Skewness of the integrated profile	Mean of the DM- SNR curve	Standard deviation of the DM-SNR curve	Excess kurtosis of the DM-SNR curve	Skewness of the DM-SNR curve
0	140.562500	55.683782	-0.234571	-0.699648	3.199833	19.110426	7.975532	74.24222
1	102.507812	58.882430	0.465318	-0.515088	1.677258	14.860146	10.576487	127.39358
2	103.015625	39.341649	0.323328	1.051164	3.121237	21.744669	7.735822	63.17190
3	136.750000	57.178449	-0.068415	-0.636238	3.642977	20.959280	6.896499	53.59366

```
1 dataset.shape
```

(17898, 9)

```
1 dataset.describe()
```

	Mean of the integrated profile	Standard deviation of the integrated profile	Excess kurtosis of the integrated profile	Skewness of the integrated profile	Mean of the DM-SNR curve	Standard deviation of the DM- SNR curve	k
<b>count</b>	17898.000000	17898.000000	17898.000000	17898.000000	17898.000000	17898.000000	17
<b>mean</b>	111.079968	46.549532	0.477857	1.770279	12.614400	26.326515	
<b>std</b>	25.652935	6.843189	1.064040	6.167913	29.472897	19.470572	
<b>min</b>	5.812500	24.772042	-1.876011	-1.791886	0.213211	7.370432	
<b>25%</b>	100.929688	42.376018	0.027098	-0.188572	1.923077	14.437332	
<b>50%</b>	115.078125	46.947479	0.223240	0.198710	2.801839	18.461316	
<b>75%</b>	127.085938	51.023202	0.473325	0.927783	5.464256	28.428104	

```
1 dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 17898 entries, 0 to 17897
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Mean of the integrated profile	17898 non-null	float64
1	Standard deviation of the integrated profile	17898 non-null	float64
2	Excess kurtosis of the integrated profile	17898 non-null	float64
3	Skewness of the integrated profile	17898 non-null	float64
4	Mean of the DM-SNR curve	17898 non-null	float64
5	Standard deviation of the DM-SNR curve	17898 non-null	float64
6	Excess kurtosis of the DM-SNR curve	17898 non-null	float64
7	Skewness of the DM-SNR curve	17898 non-null	float64
8	target_class	17898 non-null	int64

```
dtypes: float64(8), int64(1)
memory usage: 1.2 MB
```

## ▼ imputation

```
1 dataset.isnull().sum()
```

```
Mean of the integrated profile      0
Standard deviation of the integrated profile  0
Excess kurtosis of the integrated profile  0
Skewness of the integrated profile  0
Mean of the DM-SNR curve           0
Standard deviation of the DM-SNR curve  0
Excess kurtosis of the DM-SNR curve  0
Skewness of the DM-SNR curve        0
target_class                       0
dtype: int64
```

## ▼ handling column names

```
1 dataset.columns
2 # column names have spaces that can be trimmed & names can be shortened
```

```
Index([' Mean of the integrated profile',
      ' Standard deviation of the integrated profile',
      ' Excess kurtosis of the integrated profile',
      ' Skewness of the integrated profile', ' Mean of the DM-SNR curve',
      ' Standard deviation of the DM-SNR curve',
      ' Excess kurtosis of the DM-SNR curve', ' Skewness of the DM-SNR curve',
      'target_class'],
      dtype='object')
```

## ▼ stripping column names

```
1 dataset.columns = dataset.columns.str.strip()
2 # stripping column names
3 dataset.columns

Index(['Mean of the integrated profile',
      'Standard deviation of the integrated profile',
      'Excess kurtosis of the integrated profile',
      'Skewness of the integrated profile', 'Mean of the DM-SNR curve',
      'Standard deviation of the DM-SNR curve',
      'Excess kurtosis of the DM-SNR curve', 'Skewness of the DM-SNR curve',
      'target_class'],
      dtype='object')
```

## ▼ renaming column names

```
1 dataset.columns = ['Mean', 'SD', 'Kurtosis', 'Skewness', 'Mean_DM-SNR', 'SD_DM-SNR', 'Kurtosis_DM-SNR', 'Skewness_DM-SNR', 'target']
2 dataset.columns

Index(['Mean', 'SD', 'Kurtosis', 'Skewness', 'Mean_DM-SNR', 'SD_DM-SNR',
      'Kurtosis_DM-SNR', 'Skewness_DM-SNR', 'target_class'],
      dtype='object')
```

## ▼ checking dataset after handling column names

```
1 dataset.head()
```

	Mean	SD	Kurtosis	Skewness	Mean_DM-SNR	SD_DM-SNR	Kurtosis_DM-SNR	Skewness_DM-SNR
0	140.562500	55.683782	-0.234571	-0.699648	3.199833	19.110426	7.975532	74.24220
1	102.507812	58.882430	0.465318	-0.515088	1.677258	14.860146	10.576487	127.39350
2	103.015625	39.341649	0.323328	1.051164	3.121237	21.744669	7.735822	63.17190

## ▼ understanding target

```
1 dataset['target_class'].value_counts()
```

```
target_class
0    16259
1     1639
Name: count, dtype: int64
```

## ▼ EDA

```
1 import seaborn as sns
```

## ▼ hist plot for target

```
1 sns.histplot(x=dataset['target_class'])
```

<Axes: xlabel='target\_class', ylabel='Count'>




### ▼ hist plot for features

```

1 plt.figure(figsize=(12, 12))
2 plt.subplot(4, 2, 1)
3 dp1 = sns.histplot(x=dataset['Mean'])
4
5 plt.subplot(4, 2, 2)
6 dp2 = sns.histplot(x=dataset['SD'])
7
8 plt.subplot(4, 2, 3)
9 dp3 = sns.histplot(x=dataset['Kurtosis'])
10
11 plt.subplot(4, 2, 4)
12 dp4 = sns.histplot(x=dataset['Skewness'])
13
14 plt.subplot(4, 2, 5)
15 dp5 = sns.histplot(x=dataset['Mean_DM-SNR'])

```

```
16
17 plt.subplot(4, 2, 6)
18 dp6 = sns.histplot(x=dataset['SD_DM-SNR'])
19
20 plt.subplot(4, 2, 7)
21 dp7 = sns.histplot(x=dataset['Kurtosis_DM-SNR'])
22
23 plt.subplot(4, 2, 8)
24 dp8 = sns.histplot(x=dataset['Skewness_DM-SNR'])
25
26 plt.show()
```



### ▼ box plot for features




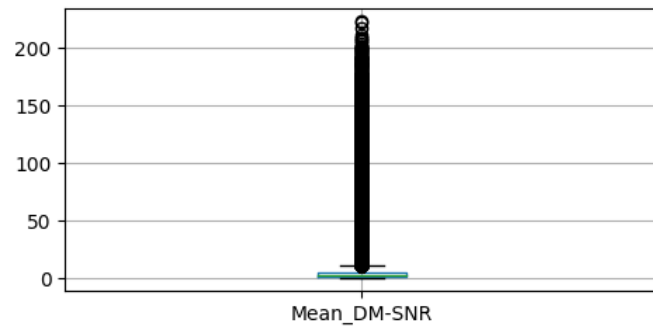
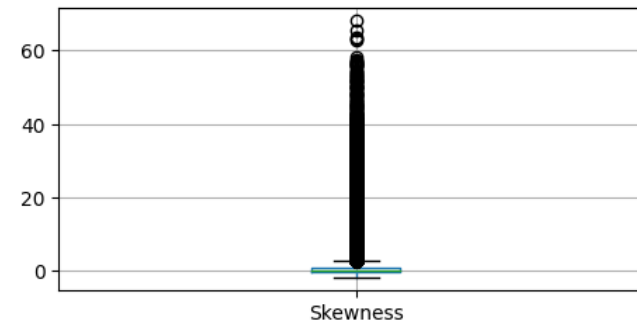
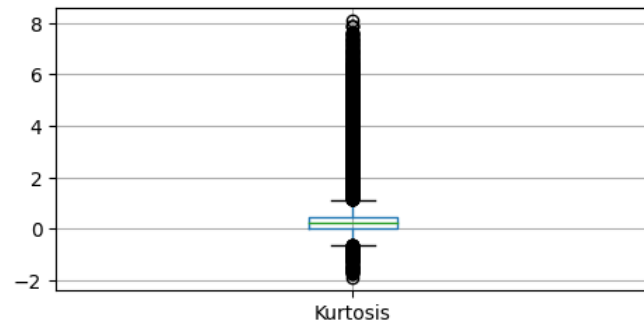


```

1 # sns.set_style(5)
2 plt.figure(figsize=(12, 12))
3 plt.subplot(4, 2, 1)
4 f1 = dataset.boxplot(column='Mean')
5
6 plt.subplot(4, 2, 2)

```



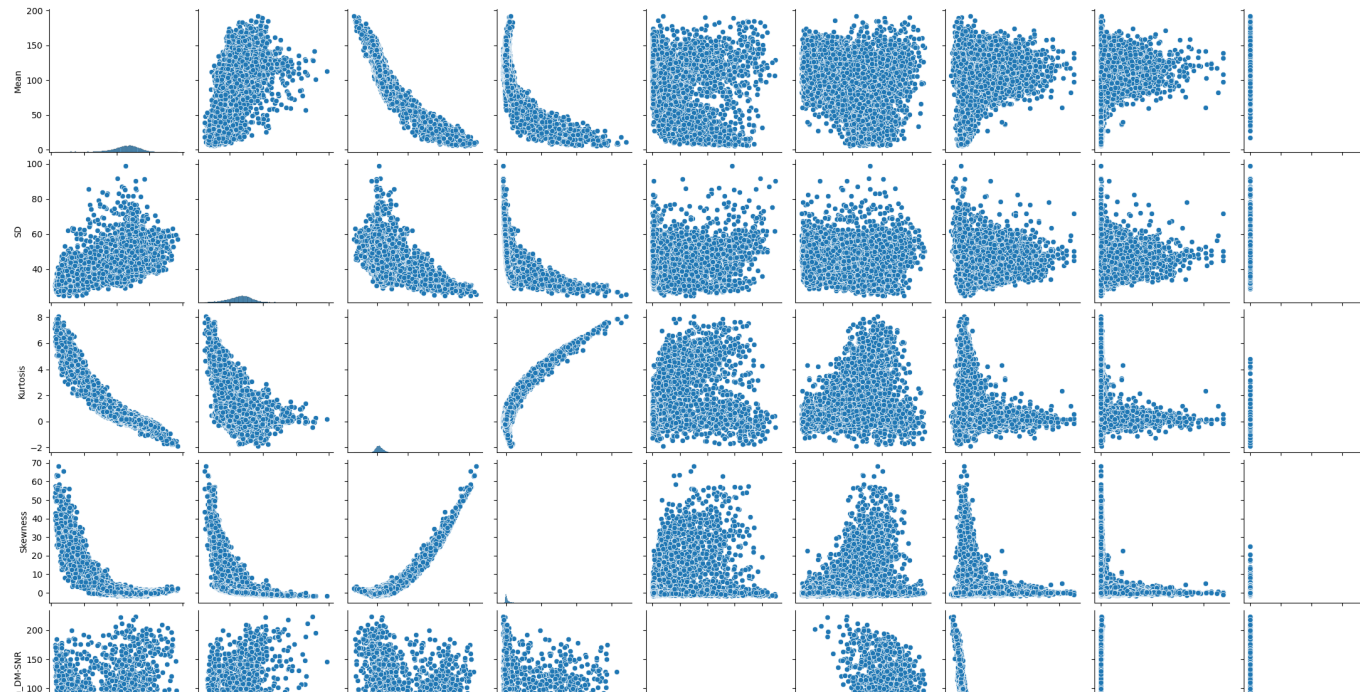
```
7 f2 = dataset.boxplot(column='SD')
8
9 plt.subplot(4, 2, 3)
10 f3 = dataset.boxplot(column='Kurtosis')
11
12 plt.subplot(4, 2, 4)
13 f4 = dataset.boxplot(column='Skewness')
14
15 plt.subplot(4, 2, 5)
16 f5 = dataset.boxplot(column='Mean_DM-SNR')
17
18 plt.subplot(4, 2, 6)
19 f6 = dataset.boxplot(column='SD_DM-SNR')
20
21 plt.subplot(4, 2, 7)
22 f7 = dataset.boxplot(column='Kurtosis_DM-SNR')
23
24 plt.subplot(4, 2, 8)
25 f8 = dataset.boxplot(column='Skewness_DM-SNR')
26
27 plt.show()
```



## ▼ Pairplot

```
1 sns.pairplot(dataset)
```

```
c:\users\surya\appdata\local\programs\python\python39\lib\site-packages\seaborn\axisgrid.p
self._figure.tight_layout(*args, **kwargs)
<seaborn.axisgrid.PairGrid at 0x1f76ae0cc10>
```



## ▼ identify X & Y



```
1 x = dataset.iloc[ : , :-1].values
2 x[:2]
```

```
array([[140.5625    ,  55.68378214, -0.23457141, -0.6996484 ,
         3.19983278,  19.11042633,  7.97553179,  74.24222492],
       [102.5078125 ,  58.88243001,  0.46531815, -0.51508791,
         1.67725752,  14.86014572, 10.57648674, 127.3935796 ]])
```



```
1 y = dataset.iloc[ : , -1].values
2 y[:5]
```

```
array([0, 0, 0, 0, 0], dtype=int64)
```

## ▼ Splitting

```
1 from sklearn.model_selection import train_test_split
```

```
1 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)
```

```
1 x_train[:2]
```

```
array([[ 1.20640625e+02,  4.78429616e+01,  2.57962577e-01,
        -9.06199080e-02,  8.04849498e+00,  3.51982345e+01,
         4.81978426e+00,  2.35283829e+01],
       [ 1.16554688e+02,  4.87029915e+01,  1.97625250e-01,
         2.32600230e-01,  3.04180602e+00,  1.66106785e+01,
         8.16618510e+00,  8.48467094e+01]])
```

```
1 y_train[:5]
```

```
array([0, 0, 1, 0, 0], dtype=int64)
```

## ▼ PreProcessing

### ▼ Feature scaling

```
1 from sklearn.preprocessing import StandardScaler
```

```
1 sc = StandardScaler()
```

```
1 x_train = sc.fit_transform(x_train)
2 x_train[:2]
```

```
array([[ 0.37710226,  0.18643192, -0.21304291, -0.30376654, -0.16236693,
         0.44549441, -0.76498842, -0.76290015],
       [ 0.21958464,  0.31196312, -0.26882602, -0.25234846, -0.32948472,
        -0.50223528, -0.02369848, -0.18468066]])
```

```
1 x_test = sc.fit_transform(x_test)
2 x_test[:2]
```

```
array([[ -0.39574205, -0.106581  , -0.2062192 , -0.23805528, -0.33389687,
        -0.60628785,  0.24052962,  0.09688056],
       [ 0.3324866  ,  0.77538991, -0.28588937, -0.3480077  , -0.29508947,
        -0.32389767, -0.10358973, -0.30403506]])
```

## ▼ Linear SVM

## ▼ Modeling - Linear SVM

```
1 from sklearn.svm import SVC
```

```
1 lsvmc = SVC(C=1, kernel='linear', random_state=0)
```

## ▼ Training - Linear SVM

```
1 lsvmc.fit(x_train, y_train)
```

```
▼ SVC
SVC(C=1, kernel='linear', random_state=0)
```

## ▼ Prediction - Linear SVM

```
1 lin_y_pred = lsvmc.predict(x_test)
2 lin_y_pred[:5]

array([0, 0, 0, 0, 0], dtype=int64)
```

## ▼ Evaluation - Linear SVM

### ▼ confusion\_matrix

```
1 from sklearn.metrics import confusion_matrix
```

```
1 confusion_matrix(y_test, lin_y_pred)
```

```
array([[3289, 17],
       [ 37, 237]], dtype=int64)
```

### ▼ classification\_report

```
1 from sklearn.metrics import classification_report
```

```
1 print(classification_report(y_test, lin_y_pred))
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	3306
1	0.93	0.86	0.90	274
accuracy			0.98	3580
macro avg	0.96	0.93	0.94	3580
weighted avg	0.98	0.98	0.98	3580

### ▼ accuracy\_score

```
1 from sklearn.metrics import accuracy_score
```

```
1 accuracy_score(y_test, lin_y_pred)
```

```
0.9849162011173185
```

### ▼ precision\_score

```
1 from sklearn.metrics import precision_score
```

```
1 precision_score(y_test, lin_y_pred)
```

```
0.9330708661417323
```

### ▼ recall\_score

```
1 from sklearn.metrics import recall_score
```



```
1 recall_score(y_test, lin_y_pred)
```

```
0.864963503649635
```

## ▼ Kernel SVM (c=1, kernel='rbf')

## ▼ Modeling - Kernel SVM (c=1, kernel='rbf')

```
1 from sklearn.svm import SVC
```

```
1 ksvmc = SVC(C=1, kernel='rbf', random_state=0)
```

## ▼ Training - Kernel SVM (c=1, kernel='rbf')

```
1 ksvmc.fit(x_train, y_train)
```

▼ SVC

SVC(C=1, random\_state=0)

## ▼ Prediction - Kernel SVM (c=1, kernel='rbf')

```
1 kernel_y_pred = ksvmc.predict(x_test)
2 kernel_y_pred[:5]
```

```
array([0, 0, 0, 0, 0], dtype=int64)
```

## ▼ Evaluation - Kernel SVM (c=1, kernel='rbf')

### ▼ confusion\_matrix

```
1 from sklearn.metrics import confusion_matrix
```

```
1 confusion_matrix(y_test, kernel_y_pred)
```

```
array([[3285,  21],
       [ 40, 234]], dtype=int64)
```

### ▼ classification\_report

```
1 from sklearn.metrics import classification_report
```

```
1 print(classification_report(y_test, kernel_y_pred))
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	3306
1	0.92	0.85	0.88	274
accuracy			0.98	3580
macro avg	0.95	0.92	0.94	3580
weighted avg	0.98	0.98	0.98	3580

### ▼ accuracy\_score

```
1 from sklearn.metrics import accuracy_score
```

```
1 accuracy_score(y_test, kernel_y_pred)
```

```
0.9829608938547486
```

### ▼ precision\_score

```
1 from sklearn.metrics import precision_score
```

```
1 precision_score(y_test, kernel_y_pred)
```

```
0.9176470588235294
```

### ▼ recall\_score

```
1 from sklearn.metrics import recall_score
```

```
1 recall_score(y_test, kernel_y_pred)
```

```
0.8540145985401459
```

### ▼ roc\_curve

```
1 from sklearn.metrics import roc_curve
```

```
1 fpr, tpr, thresholds = roc_curve(y_test, kernel_y_pred)
```

```
1 print("False Positive Rate: ", fpr)
2 print("True Positive Rate: ", tpr)
3 print("Thresholds: ", thresholds)
```

```
False Positive Rate: [0.          0.00635209 1.          ]
True Positive Rate: [0.          0.8540146 1.          ]
Thresholds: [inf  1.  0.]
```

```
1 plt.plot(fpr, tpr, linewidth=2)
2 plt.xlabel("fpr")
3 plt.ylabel("tpr")
4 plt.title("ROC curve - fpr vs tpr")
5 plt.show()
```

## ROC curve - fpr vs tpr

### ▼ roc\_auc\_score

```
1 from sklearn.metrics import roc_auc_score
```

```
1 roc_auc_score(y_test, kernel_y_pred)
2 # The higher the AUC, the better the model's performance at distinguishing
3 # between the positive and negative classes.
4 # An AUC score of 1 means the classifier can perfectly distinguish
5 # between all the Positive and the Negative class points.
```

```
0.9238312557129043
```

### ▼ Kernel SVM (c=100, kernel='poly')

### ▼ Modeling - Kernel SVM (c=100, kernel='poly')

```
1 from sklearn.svm import SVC
```

```
1 ksvm100poly = SVC(C=1, kernel='poly', random_state=0)
```

### ▼ Training - Kernel SVM (c=100, kernel='poly')

```
1 ksvm100poly.fit(x_train, y_train)
```

▼ SVC

## ▼ Prediction - Kernel SVM (c=100, kernel='poly')

```
1 y_pred_ksvm100poly = ksvm100poly.predict(x_test)
2 y_pred_ksvm100poly[:5]

array([0, 0, 0, 0, 0], dtype=int64)
```

## ▼ Evaluation - Kernel SVM (c=100, kernel='poly')

### ▼ confusion\_matrix

```
1 from sklearn.metrics import confusion_matrix
```

```
1 confusion_matrix(y_test, y_pred_ksvm100poly)
```

```
array([[3286,  20],
       [ 42, 232]], dtype=int64)
```

### ▼ classification\_report

```
1 from sklearn.metrics import classification_report
```

```
1 print(classification_report(y_test, y_pred_ksvm100poly))
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	3306

1	0.92	0.85	0.88	274
accuracy			0.98	3580
macro avg	0.95	0.92	0.94	3580
weighted avg	0.98	0.98	0.98	3580

### ▼ accuracy\_score

```
1 from sklearn.metrics import accuracy_score
```

```
1 accuracy_score(y_test, y_pred_ksvm100poly)
```

```
0.9826815642458101
```

### ▼ precision\_score

```
1 from sklearn.metrics import precision_score
```

```
1 precision_score(y_test, y_pred_ksvm100poly)
```

```
0.9206349206349206
```

### ▼ recall\_score

```
1 from sklearn.metrics import recall_score
```

```
1 recall_score(y_test, y_pred_ksvm100poly)
```

```
0.8467153284671532
```

## ▼ roc\_curve

```
1 from sklearn.metrics import roc_curve
```

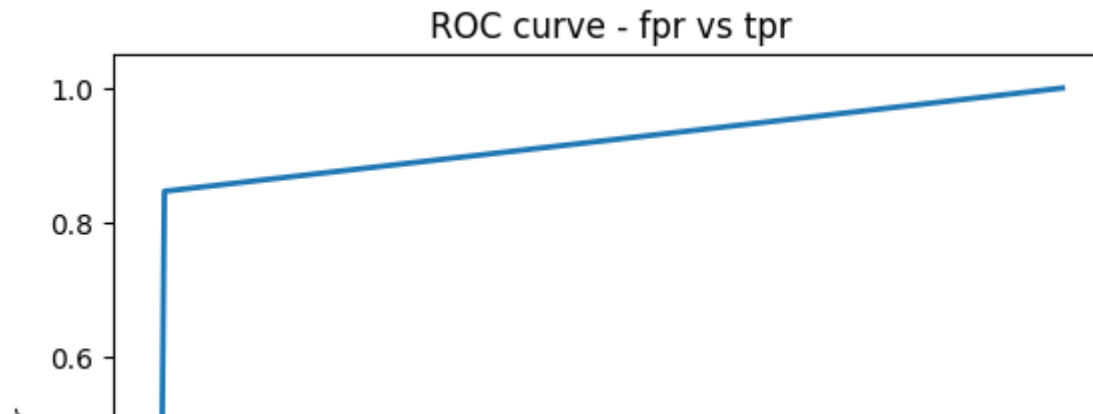
```
1 fpr, tpr, thresholds = roc_curve(y_test, y_pred_ksvm100poly)
```

```
1 print("False Positive Rate: ", fpr)
2 print("True Positive Rate: ", tpr)
3 print("Thresholds: ", thresholds)
```

```
False Positive Rate: [0.      0.00604961 1.      ]
True Positive Rate: [0.      0.84671533 1.      ]
Thresholds: [inf  1.  0.]
```

```
1 plt.plot(fpr, tpr, linewidth=2)
2 plt.xlabel("fpr")
3 plt.ylabel("tpr")
4 plt.title("ROC curve - fpr vs tpr")
5 plt.show()
```





### ▼ roc\_auc\_score

```
1 from sklearn.metrics import roc_auc_score
```

```
1 roc_auc_score(y_test, y_pred_ksvm100poly)
2 # The higher the AUC, the better the model's performance at distinguishing
3 # between the positive and negative classes.
4 # An AUC score of 1 means the classifier can perfectly distinguish
5 # between all the Positive and the Negative class points.
```

```
0.9203328608457969
```

### ▼ Kernel SVM (c=100, kernel='sigmoid')

### ▼ Modeling - Kernel SVM (c=100, kernel='sigmoid')

```
1 from sklearn.svm import SVC
```

```
1 ksvm100sigmoid = SVC(C=100, kernel='sigmoid', random_state=0)
```

### ▼ Training - Kernel SVM (c=100, kernel='sigmoid')

```
1 ksvm100sigmoid.fit(x_train, y_train)
```

```
▼ SVC
SVC(C=100, kernel='sigmoid', random_state=0)
```

### ▼ Prediction - Kernel SVM (c=100, kernel='sigmoid')

```
1 y_pred_ksvm100sigmoid = ksvm100sigmoid.predict(x_test)
2 y_pred_ksvm100sigmoid[:5]
```

```
array([0, 0, 0, 0, 0], dtype=int64)
```

### ▼ Evaluation - Kernel SVM (c=100, kernel='poly')

#### ▼ confusion\_matrix

```
1 from sklearn.metrics import confusion_matrix
```

```
1 confusion_matrix(y_test, y_pred_ksvm100sigmoid)
```

```
array([[3013, 293],
       [ 188,  86]], dtype=int64)
```

## ▼ classification\_report

```
1 from sklearn.metrics import classification_report
```

```
1 print(classification_report(y_test, y_pred_ksvm100sigmoid))
```

	precision	recall	f1-score	support
0	0.94	0.91	0.93	3306
1	0.23	0.31	0.26	274
accuracy			0.87	3580
macro avg	0.58	0.61	0.59	3580
weighted avg	0.89	0.87	0.88	3580

## ▼ accuracy\_score

```
1 from sklearn.metrics import accuracy_score
```

```
1 accuracy_score(y_test, y_pred_ksvm100sigmoid)
```

```
0.8656424581005586
```

## ▼ precision\_score

```
1 from sklearn.metrics import precision_score
```

```
1 precision_score(y_test, y_pred_ksvm100sigmoid)
```

0.22691292875989447

## ▼ recall\_score

```
1 from sklearn.metrics import recall_score
```

```
1 recall_score(y_test, y_pred_ksvm100sigmoid)
```

0.31386861313868614

## ▼ roc\_curve

- Receiver Operating Characteristic curve
- can be used to evaluate classification, mostly used for binary classifiers such as Logistic regression/classification or sigmoid classification

```
1 from sklearn.metrics import roc_curve
```

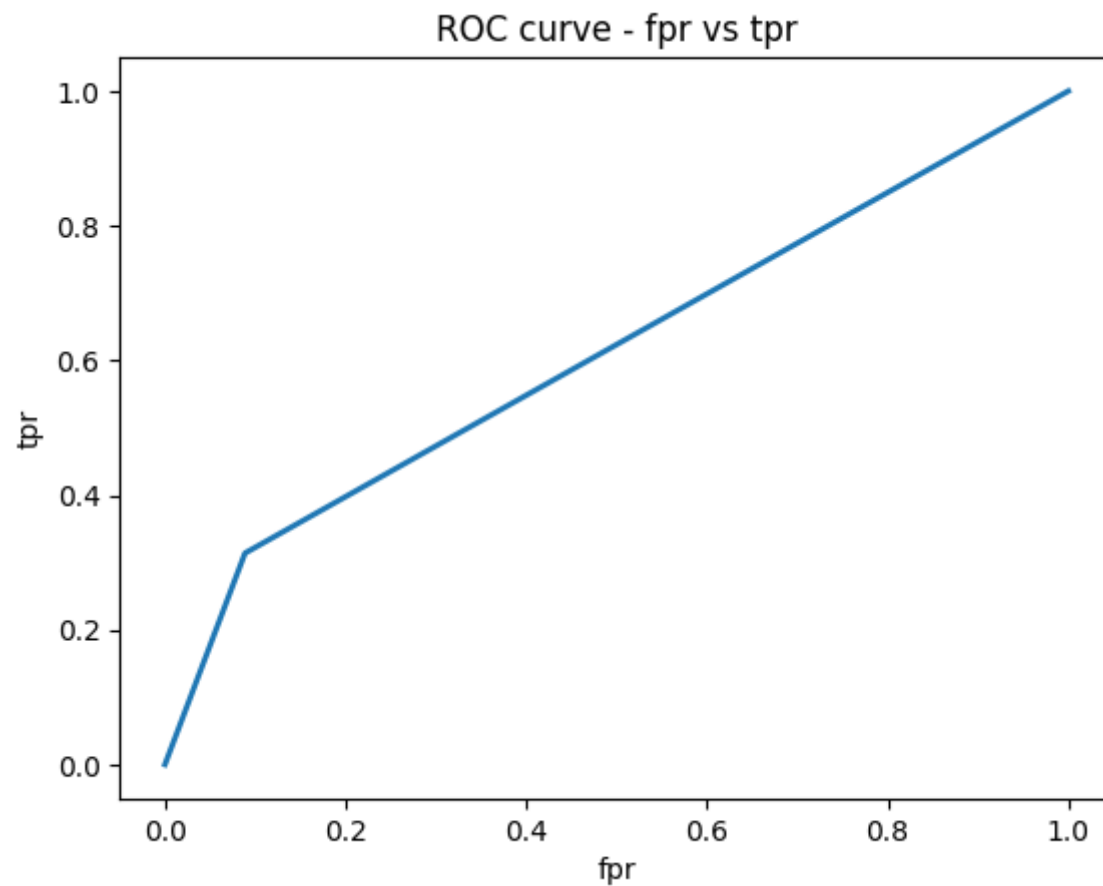
```
1 fpr, tpr, thresholds = roc_curve(y_test, y_pred_ksvm100sigmoid)
```

```
1 print("False Positive Rate: ", fpr)
2 print("True Positive Rate: ", tpr)
3 print("Thresholds: ", thresholds)
```

```
False Positive Rate: [0.          0.08862674 1.          ]
True Positive Rate: [0.          0.31386861 1.          ]
Thresholds: [inf  1.  0.]
```

```
1 plt.plot(fpr, tpr, linewidth=2)
2 plt.xlabel("fpr")
```

```
3 plt.ylabel("tpr")
4 plt.title("ROC curve - fpr vs tpr")
5 plt.show()
```



#### ▼ roc\_auc\_score

- Area Under the Receiver Operating Characteristic Curve (ROC AUC) from prediction scores.

```
1 from sklearn.metrics import roc_auc_score
```

```
1 roc_auc_score(y_test, y_pred_ksvm100sigmoid)
2 # Area Under the Receiver Operating Characteristic Curve (ROC AUC)
3 # from prediction scores.
4 # The higher the AUC, the better the model's performance at distinguishing
5 # between the positive and negative classes.
6 # An AUC score of 1 means the classifier can perfectly distinguish
7 # between all the Positive and the Negative class points.
```

0.612620936938369

